

Sean Carda
PID: A59009786
ECE 253 – Image Processing
November 7, 2021
Homework 2

Academic Integrity Policy: Integrity of scholarship is essential for an academic community. The University expects that both faculty and students will honor this principle and in so doing protect the validity of University intellectual work. For students, this means that all academic work will be done by the individual to whom it is assigned, without unauthorized aid of any kind.

By including this in my report, I agree to abide by the Academic Integrity Policy mentioned above.

Problem 1:

Code:

```
29 # Define a function for adaptive histogram equalization.
30 def AHE(im, win_size):
31     # For looping purposes, grab the dimensions of the image.
32     im = im[:, :, 0]
33     im_size = im.shape
34
35     # Also, instantiate an empty output image of the same dimensions as the input image.
36     output = np.zeros([im_size[0], im_size[1]])
37
38     # Grab the offset produced by the win_size.
39     offset = int((win_size - 1) / 2)
40
41     # Pad the image on all four sides such that the edges of the image may be operated on.
42     im_pad = np.pad(im, offset, mode='symmetric')
43
44     # Begin looping through the image.
45     print("Starting to loop...")
46     for x in range(0, im_size[0]):
47         for y in range(0, im_size[1]):
48             # Fetch the region of interest based on the current x, y coordinate.
49             #region = im_pad[(x - offset):(x + offset), (y - offset):(y + offset)]
50             region = im_pad[(x):(x + 2*offset), (y):(y + 2*offset)]
51
52             # Calculate the rank of the coordinate by summing the boolean matrix comparing
53             # im(x, y) to region.
54             rank = sum(sum(im[x, y] > region))
55
56             # Load the output with a pixel dependent on the window size and the rank counter.
57             output[x, y] = round((rank * 255) / (win_size * win_size))
58
59     # Return the equalized image.
60     return np.uint8(output)
61
62 # File path to the beach image.
63 file_path = r"beach.png"
64
65 # Begin by importing an image.
66 A = cv2.imread(file_path)
67
```

```

68 # Show the original image.
69 cv2.imshow('Beach', A)
70 cv2.waitKey(0)
71
72 # For each window size, compute the
73 for i in [33, 65, 129]:
74     # Test the function.
75     print("Equalizing the image...")
76     B = AHE(cv2.imread(file_path), i)
77     print("Equalization complete!")
78
79     # Show the equalized image.
80     cv2.imshow(('Beach_AHE_' + str(i)), B)
81     cv2.waitKey(0)
82
83     # Save the computed image.
84     file_name = 'beach_image_win_' + str(i) + '.png'
85     cv2.imwrite(file_name, B)
86
87
88 # Show the difference in performance between AHE and HE by using opencv to compute
89 # the HE image.
90 A = cv2.imread(file_path)
91 B = cv2.equalizeHist(A[:, :, 0])
92
93 cv2.imshow('Beach_AHE', B)
94 cv2.waitKey(0)
95 cv2.imwrite(('beach_image_he.png'), B)

```

Output:

Original Image:



AHE Image with a Window of Size 33:



AHE Image with a Window of Size 65:



AHE Image with a Window of Size 129:



Simple HE Image:



Questions:**a)**

While the AHE image, aesthetically, does not look as pleasing as the original due to the wide spread of pixel intensities, the AHE image reveals significantly more information about the original image. Things such as the texture of the ground, the building, clouds in the sky, and even the fact that there exists a fourth person in the building, are revealed in the AHE image. The HE image is aesthetically more pleasing than the original due to the improved spread and balance of pixel intensities but does not reveal much more information that we could not already ascertain from the original.

b)

If we consider which enhancement method improves the general appearance of the given image, the HE approach is superior due to the smoother spread of pixel intensities throughout the image. If we consider which enhancement method is superior for something like a machine to read, the AHE image is superior due to the fact that AHE image contains more meaningful information. However, it is not necessarily true that this will hold for any image; it is possible that the AHE algorithm will distort the given image depending on the window size, and therefore requires more testing to reach a concrete conclusion.

Problem 2:

Part i:

Code:

```
104 #-----
105 # i)
106 #-----
107 cv2.destroyAllWindows()
108
109 # a)
110
111 # First, load the shapes image.
112 file_path = r'circles_lines.jpg'
113 shapes = (cv2.imread(file_path))
114 shapes_en = cv2.resize(shapes, (501, 489), interpolation=cv2.INTER_AREA)
115 cv2.imwrite('shapes_enlarged.jpg', shapes_en)
116 cv2.imshow('Shapes', shapes_en)
117 cv2.waitKey(0)
118
119 # Turn the image into a binary image for morphological operations. Here, pixels with intensities
120 # above 130 are given a value of 1. Otherwise, give a value of 0. For binary purposes, divide
121 # by 255 to have values of 0 and 1.
122 shapes_binary = np.uint8(np.where(shapes[:, :, 0] > 130, 255, 0) / 255)
123 bin_shapes_enlarge = cv2.resize(shapes_binary * 255, (501, 489), interpolation=cv2.INTER_AREA) # Enlarge
124 cv2.imwrite('circles_binary.jpg', bin_shapes_enlarge)
125 cv2.imshow('Shapes_Binary', bin_shapes_enlarge)
126 cv2.waitKey(0)
127
128 # First, instantiate a structure element for the morphological operation.
129 element = np.array([[0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0],
130                    [0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0],
131                    [0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0],
132                    [0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0],
133                    [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1],
134                    [1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1],
135                    [0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0],
136                    [0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0],
137                    [0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0],
138                    [0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0]])
139
140
141 # Perform opening with the structure element defined above.
142 circles = morphology.opening(shapes_binary, element)
143
144 # Show the result.
145 circles_en = cv2.resize(circles * 255, (501, 489), interpolation=cv2.INTER_AREA)
146 cv2.imwrite('circles_opened.jpg', circles_en)
147 cv2.imshow('Circles', circles_en)
148 cv2.waitKey(0)
149
150 # b)
151
152 # Now, label the circles in the morphed image.
153 labeled_circles = ndimage.measurements.label(circles)
154
155 # Plot the labeled objects.
156 plt.imshow(labeled_circles[0])
157 plt.colorbar()
158 plt.xlabel('Columns')
159 plt.ylabel('Rows')
160 plt.title('Circles Colored by Connected Component Analysis')
161 plt.show()
162
163 # c)
164
165 # Method for computing the centroids and areas of a list of given shapes.
166 def compute_object_stats(object_original, object_matrix, object_count, quantity):
167
168     # Instantiate a dictionary in which the centroid coordinates and areas will be stores.
169     object_dictionary = {}
170
171     # Grab the size of the input object matrix.
172     object_matrix_size = object_matrix.shape
173
174     # For every object, calculate its centroid and area.
175     for i in range(1, object_count + 1):
176         # First, calculate the area of the object.
177         valid_elements = (object_matrix == i)
178
179         # If the given quantity command is length, then compute the lengths of the objects.
180         if quantity == 'length':
181             max_length = 0
182             for c in range(0, object_matrix_size[1]):
183                 length = sum(valid_elements[:, c])
184                 if length > max_length:
185                     max_length = length
186             measure = max_length
187
188         # Otherwise, calculate the areas of the objects.
189         else:
190             measure = sum(sum(valid_elements))
191
192     # Now, find all the x and y coordinates where the matrix has valid elements.
193     # then, calculate the mean of those coordinates.
194     r_list = []
195     c_list = []
196
```



```

197     # Locate valid coordinates.
198     for r in range(0, object_matrix_size[0]):
199         for c in range(0, object_matrix_size[1]):
200             if valid_elements[r, c] > 0:
201                 r_list.append(r)
202                 c_list.append(c)
203
204     # Take their mean.
205     r_mean = round(mean(r_list))
206     c_mean = round(mean(c_list))
207
208     # Load the area and centroid as a new value in the dictionary.
209     object_dictionary['Object ' + str(i)] = [measure, [r_mean, c_mean]]
210
211 # Create a new image and verify that the centroids have been calculated correctly.
212 centroid_im = np.zeros([object_matrix_size[0], object_matrix_size[1], 3])
213 centroid_im[:, :, 0] = object_original
214 centroid_im[:, :, 1] = object_original
215 centroid_im[:, :, 2] = object_original
216
217 # For all the objects in the dictionary, grab the calculated coordinates and set the values
218 # of the original image at those coordinates to blue.
219 for key in object_dictionary:
220     coordinates = (object_dictionary.get(key))[1]
221     for r in range(0, object_matrix_size[0]):
222         for c in range(0, object_matrix_size[1]):
223             if r == coordinates[0] and c == coordinates[1]:
224                 centroid_im[r, c, 1] = 0
225                 centroid_im[r, c, 2] = 0
226
227 # Return both the centroid image as well as the object dictionary containing the relevant
228 # info on all of the objects in the image.
229 return [centroid_im, object_dictionary]
230
231
232 # Method for printing the computed shape stats.
233 def print_stats(object_dict, measure):
234     print("\n--Object--\t--" + measure + "--\t--Centroid [row, col]--")
235     for key in object_dict:
236         data = object_dict.get(key)
237         print(key + ": \t" + str(data[0]) + "\t\t" + str(data[1]))
238
239 # Compute the object stats for the circle objects.
240 circle_stats = compute_object_stats(circles, labeled_circles[0], labeled_circles[1], 'area')
241
242 # Print the stats computed for the circles.
243 print_stats(circle_stats[1], 'Area')
244
245 # Display the positions of the centroids to verify that the centroids match the positions of the shapes.
246 centroid_en = cv2.resize(255 * circle_stats[0], (501, 489), interpolation=cv2.INTER_AREA) # Enlarge.
247 cv2.imwrite('circle_centroids.jpg', centroid_en)
248 cv2.imshow('Centroids for Circles', centroid_en)
249 cv2.waitKey(0)

```

Output:

Original Image:



Binary Image:





Image After Opening:



Structuring Element Used:

For this problem, the structuring element was a disk of radius 5. The matrix itself was a 10 x 10 matrix, with ones in places where pixels in the binary image should be kept, and zeros where pixels in the binary image should be removed. This structuring element was successful in removing most of the lines in the image. However, some of the preserved circles have areas larger than others due to the fact that some circles were intersecting with lines.

Image After Connected Component Analysis (with color-bar):

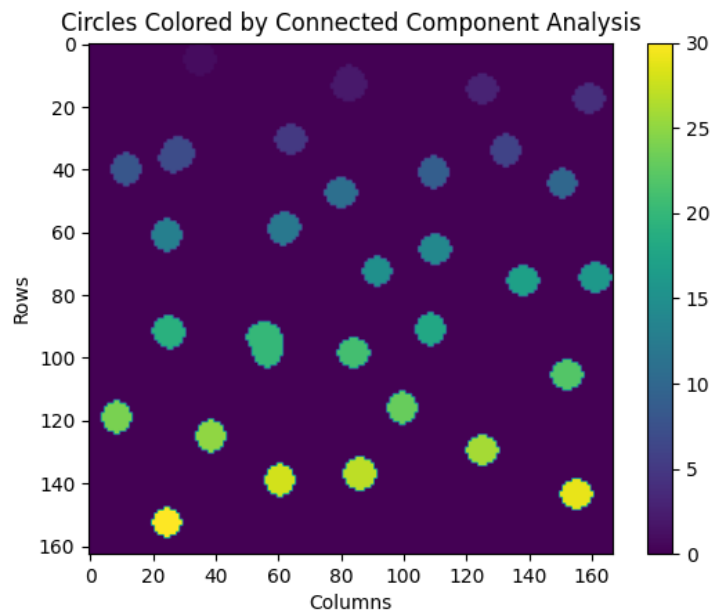


Image After Calculating Centroids:




Table of Centroid and Area Values:

--Object--	--Area--	--Centroid [row, col]--
Object 1:	89	[5, 35]
Object 2:	102	[13, 83]
Object 3:	78	[14, 125]
Object 4:	78	[18, 159]
Object 5:	78	[30, 64]
Object 6:	78	[34, 132]
Object 7:	100	[35, 28]
Object 8:	78	[40, 12]
Object 9:	78	[41, 110]
Object 10:	68	[44, 150]
Object 11:	78	[48, 80]
Object 12:	85	[59, 62]
Object 13:	78	[61, 24]
Object 14:	78	[66, 110]
Object 15:	68	[72, 92]
Object 16:	78	[74, 161]
Object 17:	78	[76, 138]
Object 18:	78	[91, 108]
Object 19:	85	[92, 25]
Object 20:	132	[96, 56]
Object 21:	78	[98, 84]
Object 22:	78	[106, 152]
Object 23:	78	[116, 100]
Object 24:	78	[119, 8]
Object 25:	78	[125, 38]
Object 26:	78	[130, 125]
Object 27:	89	[137, 86]
Object 28:	78	[139, 60]
Object 29:	78	[144, 155]
Object 30:	68	[152, 24]

Part ii:

Code:

```
251 #-----
252 # ii)
253 #-----
254 cv2.destroyAllWindows()
255
256 # First, load the shapes image.
257 file_path = r'lines.jpg'
258 shapes = (cv2.imread(file_path))
259 shapes_en = cv2.resize(shapes, (462, 552), interpolation=cv2.INTER_AREA)
260 cv2.imwrite('shapes_enlarged.jpg', shapes_en)
261 cv2.imshow('Original Lines Image', shapes_en)
262 cv2.waitKey(0)
263
264 # Turn the image into a binary image for morphological operations. Here, pixels with intensities
265 # above 130 are given a value of 1. Otherwise, give a value of 0. For binary purposes, divide
266 # by 255 to have values of 0 and 1.
267 shapes_binary = np.uint8(np.where(shapes[:, :, 0] > 130, 255, 0) / 255)
268 bin_shapes_enlarge = cv2.resize(shapes_binary * 255, (462, 552), interpolation=cv2.INTER_AREA) # Enlarge
269 cv2.imwrite('binary_lines.jpg', bin_shapes_enlarge)
270 cv2.imshow('Binary Lines', bin_shapes_enlarge)
271 cv2.waitKey(0)
272
273 # a)
274
275 # First, instantiate a structure element for the morphological operation.
276 element = np.ones([13, 1])
277
278 #lines = morphology.opening(shapes_binary_filter, element)
279 lines = morphology.opening(shapes_binary, element)
280
281 # Show the result.
282 lines_en = cv2.resize(lines * 255, (462, 552), interpolation=cv2.INTER_AREA)
283 cv2.imwrite('lines_en.jpg', lines_en)
284 cv2.imshow('Lines', lines_en)
285 cv2.waitKey(0)
286
287 # b)
288
289 # Now, label the circles in the morphed image.
290 labeled_lines = ndimage.measurements.label(lines)
291
292 # Plot the labeled objects.
293 plt.imshow(labeled_lines[0])
294 plt.colorbar()
295 plt.xlabel('Columns')
296 plt.ylabel('Rows')
```

```

297 plt.title('Lines Colored by Connected Component Analysis')
298 plt.show()
299
300 # c)
301
302 # Compute the object stats for the vertical line objects.
303 line_stats = compute_object_stats(lines, labeled_lines[0], labeled_lines[1], 'length')
304
305 # Print the stats computed for the lines.
306 print_stats(line_stats[1], 'Length')
307
308 # Display the positions of the centroids to verify that the centroids match the positions of the shapes.
309 centroid_en = cv2.resize(255 * line_stats[0], (462, 552), interpolation=cv2.INTER_AREA) # Enlarge.
310 cv2.imwrite('centroids_lines.jpg', centroid_en)
311 cv2.imshow('Centroids for Lines', centroid_en)
312 cv2.waitKey(0)

```

Output:

Original Image:



Binary Image:




Image After Opening:



Structuring Element Used:

For this problem, the structuring element was a rectangle of length 13, and width 1. The matrix itself was a 13 x 1 matrix of all ones. No zeros were used in this structuring element since they would have been redundant – only the ones needed to fit in the image. This structuring element was very efficient in removing the horizontal and diagonal lines from the image.

Image After Connected Component Analysis (with color-bar):

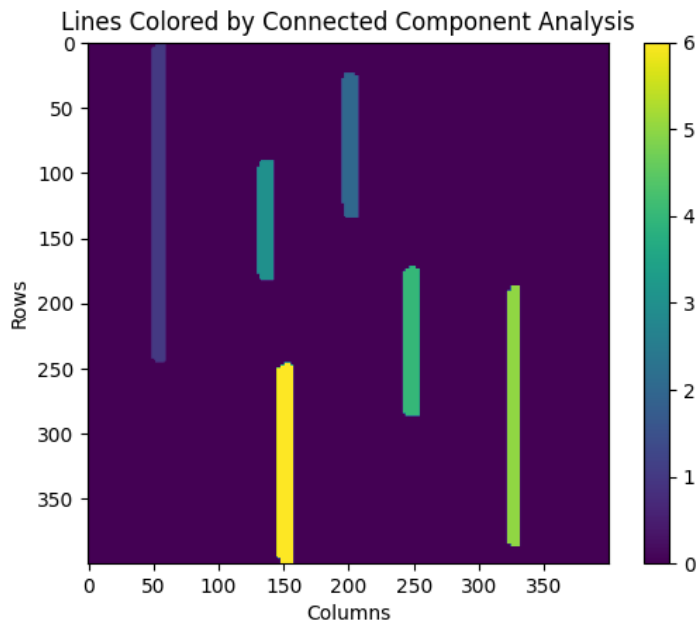


Image After Calculating Centroids:




Table of Centroid and Area Values:

--Object--	--Length--	--Centroid [row, col]--
Object 1:	113	[56, 21]
Object 2:	51	[36, 77]
Object 3:	42	[62, 52]
Object 4:	53	[105, 95]
Object 5:	92	[132, 126]
Object 6:	71	[148, 58]

Problem 3:

Part i:

Code:

Part ii:

Code:


```
357 #-----
358 # 11)
359 #-----
360
361 # This function accepts an image and a bit value and computes the Lloyd-Max quantized image
362 # for the corresponding bit value.
363 def compute_lloyd_quant(im, s):
364     # Flatten the lena image for the lloyds function to take as an input.
365     dim = im.shape
366     new_dim = [dim[0] * dim[1], 1]
367     im_flat = np.reshape(im, new_dim)
368
369     # Using the provided lloyd_python script, compute the necessary intensity partitions
370     # and corresponding transformed intensity values.
371     partition, codebook = lloyds(im_flat, [2**s])
372
373     # For this new range of intensities, quantize the original image.
374     output = np.zeros([dim[0], dim[1]])
375     for r in range(0, dim[0]):
376         for c in range(0, dim[1]):
377             # Grab the current pixel value.
378             pixel = im[r, c]
379
380             # Setup a flag marker for when to stop searching for new values.
381             flag = 0
382
383             # Instantiate an index variable to keep track of the partition index.
384             index = 0
385
386             # While we have not reached the correct transform.
387             while flag == 0:
388                 # If the current pixel value is less than the current partition bin, find all the
389                 # valid codebook values and find the value in the resulting list that is the max.
390                 if pixel < partition[index]:
391                     vals = np.where(codebook < partition[index], codebook, 0)
392                     output[r, c] = round(codebook[np.argmax(vals)])
393                     flag = 1
394
395                 # If the current pixel value is greater than all partition bins, then we want to
396                 # transform that pixel into the largest codebook intensity.
397                 elif pixel > partition[len(partition) - 1]:
398                     output[r, c] = round(codebook[2**s - 1])
399                     flag = 1
400
401                 # Adjust the index in case the conditions above have not been met.
402                 index += 1
403
404     # Return the quantized image.
405     return np.uint8(output)
406
407
408 # Now, compute the mean squared error between the original image and the uniform quantized images
409 # and Lloyd-Max quantized images for bits 1 through 7. To do this, we will create a method
410 # that accepts an image and a quantization method and computes the MSE between the original image
411 # and the quantized image.
412 def compute_mse(im, quant_method):
413     # Instantiate a dictionary to keep track of the MSE relative to the bit number.
414     mse_vals = {}
415
416     # Compute the quantized image for every bit level. Instantiate an empty image set for
417     # the quantization images.
418     size = im.shape
419     quantized_image_set = {}
420
421     # Depending on the requested quantization method, compute the quantization images for the
422     # 7 different bits.
423     if quant_method == 'Lloyd':
424         for i in range(1, 8):
425             print('Computing Lloyd quantization for bit ' + str(i) + '...')
426             quantized_image_set[i] = compute_lloyd_quant(im, i)
427     else:
428         for i in range(1, 8):
429             print('Computing Uniform quantization for bit ' + str(i) + '...')
430             quantized_image_set[i] = compute_uniform_quant(im, i)
431
432     # For every quantized image, compute the MSE values.
433     print('Computing MSE values...')
434     for i in range(0, 7):
435         # Grab the current quantized image.
436         quant_image = np.uint64(quantized_image_set[i + 1])
437         original_im = np.uint64(im)
438
439         # Compute the error.
440         error = (original_im - quant_image)
441
442         error_squared = error**2
443
444         # Sum the square of every term.
445         squared_sum = sum(sum(error_squared))
446
447         # Take the mean.
448         mse = squared_sum / (size[0] * size[1])
```

```

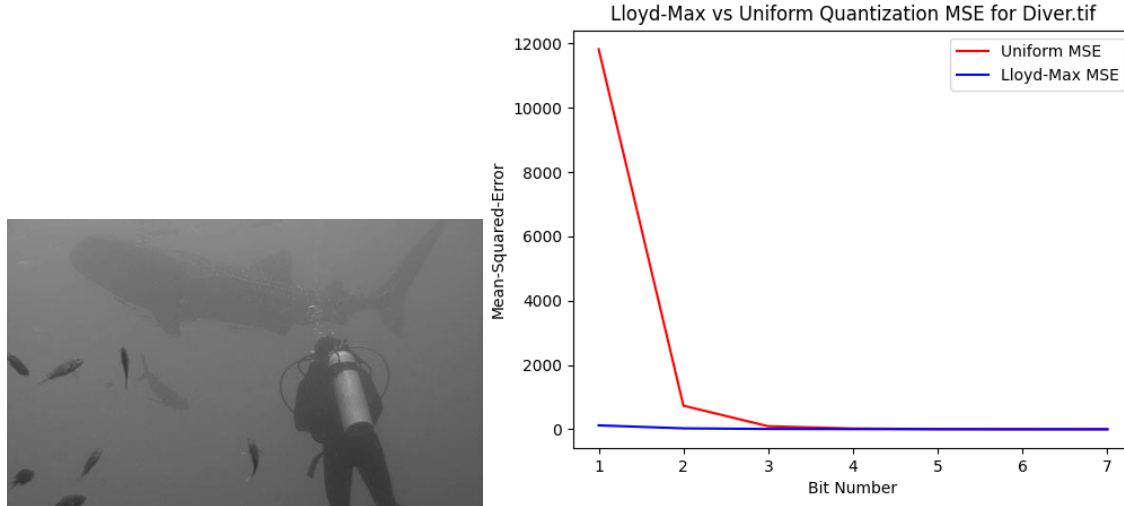
449     # Store in the corresponding bin.
450     mse_vals[i] = mse
451
452     # Return the values.
453     print('--Computation for ' + quant_method + ' quantization complete!--')
454     return mse_vals
455
456
457 # First, load the Lena512 image.
458 file_path = r'lena512.tif'
459 lena = (cv2.imread(file_path))
460
461 # Compute the MSE for the Lena image for both uniform and Lloyd quantization.
462 bit_numbers = [1, 2, 3, 4, 5, 6, 7]
463 lena_mse_unif = compute_mse(lena[:, :, 0], 'Uniform')
464 lena_mse_lloyd = compute_mse(lena[:, :, 0], 'Lloyd')
465
466 # Show the original image and plot the MSE values for both quantizers.
467 cv2.imshow('Lena Image', lena)
468 cv2.waitKey(0)
469 plt.plot(bit_numbers, lena_mse_unif.values(), color='red', label='Uniform MSE')
470 plt.plot(bit_numbers, lena_mse_lloyd.values(), color='blue', label='Lloyd-Max MSE')
471 plt.legend(loc='upper right')
472 plt.xlabel('Bit Number')
473 plt.ylabel('Mean-Squared-Error')
474 plt.title('Lloyd-Max vs Uniform Quantization MSE for Lena.tif')
475 plt.show()
476
477 # Second, load the Diver image.
478 file_path = r'diver.tif'
479 diver = (cv2.imread(file_path))
480
481 # Compute the MSE for the Lena image for both uniform and Lloyd quantization.
482 diver_mse_unif = compute_mse(diver[:, :, 0], 'Uniform')
483 diver_mse_lloyd = compute_mse(diver[:, :, 0], 'Lloyd')
484
485 # Show the original image and plot the MSE values for both quantizers.
486 cv2.imshow('Diver Image', diver)
487 cv2.waitKey(0)
488 plt.plot(bit_numbers, diver_mse_unif.values(), color='red', label='Uniform MSE')
489 plt.plot(bit_numbers, diver_mse_lloyd.values(), color='blue', label='Lloyd-Max MSE')
490 plt.legend(loc='upper right')
491 plt.xlabel('Bit Number')
492 plt.ylabel('Mean-Squared-Error')
493 plt.title('Lloyd-Max vs Uniform Quantization MSE for Diver.tif')
494

```

Output:
Original Lena Image and MSE Plot:



Original Diver Image and MSE Plot:



A likely reason for the stark gap in MSE between the uniform quantization and Lloyd-Max quantization methods, and perhaps a primary reason for why the Lloyd-Max quantizer outperforms the uniform quantizer, is due to the fact that the Lloyd-Max quantizer establishes intensity levels that are closer to the average intensity level of the original image than the uniform quantizer. This means that the gaps between the Lloyd-Max intensity levels will be smaller than the gap between the uniform intensity levels, resulting in lower error. This is the reason why the uniform quantizer produces such a large 1-bit MSE value – is that the 1-bit quantizer produces intensity levels at the most extreme ends of the pixel intensity spectrum (on a 8-bit scale, that is).

As the bits increase, the gap decreases between the two quantizers since the levels produced by both start to become closer to the levels in the 8-bit image, therefore reducing the error.


Part iii:

Code:


```
497 #-----
498 # iii)
499 #-----
500
501 # Now, compute the MSE after equalizing the histograms
502 lena_eq = cv2.equalizeHist(lena[:, :, 0])
503
504 # Compute the MSE for the equalized Lena image for both uniform and Lloyd quantization.
505 bit_numbers = [1, 2, 3, 4, 5, 6, 7]
506 lena_eq_mse_unif = compute_mse(lena_eq, 'Uniform')
507 lena_eq_mse_lloyd = compute_mse(lena_eq, 'Lloyd')
508
509 # Show the original image and plot the MSE values for both quantizers.
510 cv2.imwrite('lena_eq.jpg', lena_eq)
511 cv2.imshow('Equalized Lena Image', lena_eq)
512 cv2.waitKey(0)
513 plt.plot(bit_numbers, lena_eq_mse_unif.values(), color='red', label='Uniform MSE')
514 plt.plot(bit_numbers, lena_eq_mse_lloyd.values(), color='blue', label='Lloyd-Max MSE')
515 plt.legend(loc='upper right')
516 plt.xlabel('Bit Number')
517 plt.ylabel('Mean-Squared-Error')
518 plt.title('Lloyd-Max vs Uniform Quantization MSE for Equalized Lena.tif')
519 plt.show()
520
521 # Equalize the diver image.
522 diver_eq = cv2.equalizeHist(diver[:, :, 0])
523
524 # Compute the MSE for the Lena image for both uniform and Lloyd quantization.
525 diver_eq_mse_unif = compute_mse(diver_eq, 'Uniform')
526 diver_eq_mse_lloyd = compute_mse(diver_eq, 'Lloyd')
527
528 # Show the original image and plot the MSE values for both quantizers.
529 cv2.imwrite('diver_eq.jpg', diver_eq)
530 cv2.imshow('Equalized Diver Image', diver_eq)
531 cv2.waitKey(0)
532 plt.plot(bit_numbers, diver_eq_mse_unif.values(), color='red', label='Uniform MSE')
533 plt.plot(bit_numbers, diver_eq_mse_lloyd.values(), color='blue', label='Lloyd-Max MSE')
534 plt.legend(loc='upper right')
535 plt.xlabel('Bit Number')
536 plt.ylabel('Mean-Squared-Error')
537 plt.title('Lloyd-Max vs Uniform Quantization MSE for Equalized Diver.tif')
538 plt.show()
```

Output:

Equalized Lena Image and MSE Plots:



Equalized Diver Image and MSE Plots:



Here, after equalization, the gap between the two quantizers has moderately decreased. It is clear that the uniform quantizer performs better after an image has been equalized, while the Lloyd-Max quantizer performs slightly weaker than when the images were not equalized.

A probable explanation for this could be that due to the increased contrast in the images after equalization. The increased contrast would benefit the uniform quantizer since there would exist more pixels on the lower ends of the 8-bit intensity spectrum, thus reducing the error between regions closer to intensities 0 and 255. This would not benefit the Lloyd-Max quantizer since the number of pixels closer to the Lloyd-Max intensity levels would decrease, and thus, the error would increase due to the distribution of intensities further away from the Lloyd-Max intensities.

Part iv:

I believe that the MSE of the Lloyd-Max quantizer stays close to zero for 7-bit quantization due to the fact that by nature, the Lloyd-Max quantizer locates the optimal threshold locations for a given image. The thresholds given by the 7-bit quantizer remain close to the concentrations of intensities resulting from the equalization since the quantizer is actively seeking those points.

Problem 4:

Part i:


Code:

```
547 #-----
548 # i)
549 #-----
550
551 # Method which accepts an image and computes the uniform quantization for 10 intensity levels.
552 def compute_uniform_color_quant(im):
553     # Compute the number of intensity values.
554     levels = 10
555
556     # Compute the new intensity values for the original pixel values.
557     new_values = np.round((255 / (levels - 1)) * np.arange(levels))
558
559     # Grab the size of the image.
560     im_size = im.shape
561     print(im_size)
562
563     # Split the image into the three color bands and perform the quantization separately
564     # for each color layer.
565     print('Computing uniform quantized values...')
566     output = np.zeros([im_size[0], im_size[1], im_size[2]])
567     for r in range(0, im_size[0]):
568         for c in range(0, im_size[1]):
569             # Calculate the difference in intensities between the current pixel value
570             # and the new calculated intensities.
571             difference_blue = abs(new_values - im[r, c, 0])
572             difference_green = abs(new_values - im[r, c, 1])
573             difference_red = abs(new_values - im[r, c, 2])
574
575             # Find the minimum mapping from the old pixel value to the new intensity value.
576             out_blue = new_values[np.argmin(difference_blue)]
577             out_green = new_values[np.argmin(difference_green)]
578             out_red = new_values[np.argmin(difference_red)]
579
580             # Concatenate these values into the new output pixel.
581             output[r, c, :] = [out_blue, out_green, out_red]
582
583             if (r + 1) % 100 == 0:
584                 percent = round(100 * (1 + r) / im_size[0])
585                 print('--Computation Progress: ' + (str(percent) + '%'))
586
587     # Return the final output image.
588     return np.uint8(output)
589
590 # Load the geisel image.
591 file_path = r'geisel.jpg'
592 geisel = (cv2.imread(file_path))
593 cv2.imshow('Geisel Original', cv2.resize(geisel[476:576, 468:568], (400, 400), interpolation=cv2.INTER_AREA))
594 cv2.waitKey(0)
595
596 # Compute the uniform color quantization of the Geisel image.
597 geisel_uniform_quant = compute_uniform_color_quant(geisel)
598 cv2.imwrite('geisel_uniform_quant.jpg', geisel_uniform_quant)
599
600 # Show the resized image.
601 geisel_size = geisel.shape
602 geisel_uniform_quant_en = cv2.resize(geisel_uniform_quant[476:576, 468:568], (400, 400), interpolation=cv2.INTER_AREA)
603 cv2.imwrite('geisel_uniform_quant_en.jpg', geisel_uniform_quant_en)
604 cv2.imshow('Geisel Quantized Enhanced', geisel_uniform_quant_en)
605 cv2.imshow('Geisel Quantized', geisel_uniform_quant)
606 cv2.waitKey(0)
```

Output:
Original Image:



Zoomed-in Original:



Quantized Image:



Zoomed-in Quantized Image:



Part ii:


Code:

```
608 #-----
609 # ii)
610 #-----
611
612 # Function which accepts an RGB pixel and computes a new pixel with the closest intensities
613 # for 10 new intensity levels.
614 def find_closest_palette_color(pixel):
615     # Compute a new pixel value given that we are quantizing to 10 intensity levels.
616     new_values = np.round((255 / 9) * np.arange(10))
617
618     # Compute new intensity values for all three color layers.
619     red = new_values[np.argmin(abs(new_values - pixel[2]))]
620     green = new_values[np.argmin(abs(new_values - pixel[1]))]
621     blue = new_values[np.argmin(abs(new_values - pixel[0]))]
622
623     # Concatenate the new pixel values.
624     return [blue, green, red]
625
626 # Function which computes Floyd-Steinberg dithering quantization for a given RGB image. This
627 # function will compute the quantized image for 10 intensity levels.
628 def compute_floyd_steinberg_dither(im):
629     # First, grab the dimensions of the image.
630     size = im.shape
631
632     # Now, pad the original image due to boundary computations in the F-S algorithm. This padding
633     # only needs to be done for 1 pixel.
634     im_pad = np.pad(im, ((1, 1), (1, 1), (0, 0)), mode='symmetric')
635
636     # Instantiate an output image to return.
637     output = np.int64(np.copy(im_pad))
638
639     # Perform the algorithm.
640     print('Computing Floyd-Steinberg quantized values...')
641     for r in range(0, size[0]):
642         for c in range(0, size[1]):
643             # Grab the current pixel value from the given image.
644             old_pixel = output[r, c, :]
645
646             # Grab the new pixel based on the 'find_closest_palette_color' function.
647             new_pixel = find_closest_palette_color(old_pixel)
648
649             # Compute the error between the new pixel and the old pixel.
650             error = old_pixel - new_pixel
651
652             # Store the new pixel in the output image.
653             output[r, c, :] = new_pixel
654
655             # Now, compute the pixels surrounding the pixel we just stored.
656             output[r + 1, c, :] = np.round(output[r + 1, c, :] + ((7 / 16) * error))
657             output[r - 1, c + 1, :] = np.round(output[r - 1, c + 1, :] + ((3 / 16) * error))
658             output[r, c + 1, :] = np.round(output[r, c + 1, :] + ((5 / 16) * error))
659             output[r + 1, c + 1, :] = np.round(output[r + 1, c + 1, :] + ((1 / 16) * error))
660
661             # Check progress of computation.
662             if (r + 1) % 100 == 0:
663                 percent = round(100 * (1 + r) / size[0])
664                 print('--Computation Progress: ' + (str(percent) + '%'))
665
666     return np.uint8(output)
667
668
669 # Calculate the Floyd-Steinberg dithering quantization for the Geisel image.
670 geisel_fs_quant = compute_floyd_steinberg_dither(geisel)
671 cv2.imwrite('geisel_fs_quant.jpg', geisel_fs_quant)
672
673 # Show the quantized image.
674 geisel_fs_quant_en = cv2.resize(geisel_fs_quant[476:576, 468:568], (400, 400), interpolation=cv2.INTER_AREA)
675 cv2.imwrite('geisel_fs_quant_en.jpg', geisel_fs_quant_en)
676 cv2.imshow('Geisel Dithered Enhanced', geisel_fs_quant_en)
677 cv2.imshow('Geisel Dithered', geisel_fs_quant)
678 cv2.waitKey(0)
```


Output:
Dithered Image:



Zoomed-in Dithered Image:



Questions:**1)**

The biggest difference between the two images is that the uniformly quantized image possesses far less detail than the dithered image. And while it cannot be gleaned from far away, the dithered image has a “tiled” texture to it, where the pixels are weaved together to give the illusion of more intensity values than there really are, and thus “smoothens” the image. The texture of the quantized image is not as smooth, and the jumps between the intensity values are far more prominent due to the restrictions on the intensity levels. It should be noted that in this case, the colors on the dithering image look slightly stranger.

2)

The primary reason behind these differences is that the dithering algorithm propagates the error between the original and quantized pixels throughout the image. This is what leads to the “tiled” effect stated previously. The tiling gives the illusion that there are more intensity values than exist given the number of desired layers. This does not actually mean that there are an increased number of intensity layers than are desired.

The difference in terms of color is likely due to an error in propagating the pixel error throughout the image. As a result, some of the pixels appear distorted, and do not reflect the original colors of the image. This is likely not a flaw in the algorithm itself, but a flaw in managing the overflow of datatypes.