

Sean Carda
PID: A59009786
ECE 253 – Image Processing
November 21, 2021
Homework 3

Academic Integrity Policy: Integrity of scholarship is essential for an academic community. The University expects that both faculty and students will honor this principle and in so doing protect the validity of University intellectual work. For students, this means that all academic work will be done by the individual to whom it is assigned, without unauthorized aid of any kind.

By including this in my report, I agree to abide by the Academic Integrity Policy mentioned above.

Problem 1:

Code:

```
25 # Define a function that accepts as inputs a grayscale image and a threshold and returns an
26 # image containing the edges of the original.
27 def compute_canny_edge(im, threshold):
28     # Start the smoothing process.
29     print('-----')
30     print('---Smoothing---')
31
32     # Grab the dimensions of the image.
33     size = im.shape
34
35     # The first step is to smooth the image using the following gaussian kernel.
36     K = (1 / 159) * np.array([[2, 4, 5, 4, 2],
37                               [4, 9, 12, 9, 4],
38                               [5, 12, 15, 12, 5],
39                               [4, 9, 12, 9, 4],
40                               [2, 4, 5, 4, 2]])
41
42
43
44     # Compute the convolution between the kernel and the padded image.
45     print('Beginning the convolution for smoothing...')
46     im_smooth = compute_convolution(im, K)
47     im_smooth = ndimage.convolve(im, K)
48
49     # Start the gradient calculation process.
50     print('-----')
51     print('---Gradients---')
52
53     # Instantiate the necessary kernels to calculate the respective gradients.
54     k_x = np.array([[ -1, 0, 1],
55                    [ -2, 0, 2],
56                    [ -1, 0, 1]])
57
58     k_y = np.array([[ -1, -2, -1],
59                    [ 0, 0, 0],
60                    [ 1, 2, 1]])
61
62     # Calculate the gradients in the x and y directions.
63     print('Computing x-gradient...')
64     G_x = compute_convolution(np.float64(im_smooth), k_x)
65
66     print('Computing y-gradient...')
67     G_y = compute_convolution(np.float64(im_smooth), k_y)
68
69     # Calculate the magnitude and direction of the gradient.
70     G_mag = np.sqrt(G_x**2 + G_y**2)
71     G_dir = np.degrees(np.arctan2(G_y, G_x))
```

```

72
73 # Start the suppression process.
74 print('-----')
75 print('---NMS---')
76
77 # First, correct the degrees
78 print('Correcting the degrees matrix...')
79 G_dir_round = np.zeros([size[0], size[1]])
80 for r in range(0, size[0]):
81     for c in range(0, size[1]):
82         G_dir_round[r, c] = round_degrees(G_dir[r, c])
83
84 # Pad the gradient magnitude matrix for NMS checks.
85 G_mag_pad = np.pad(G_mag, 1, mode='symmetric')
86
87 # For every pixel in the padded image, determine whether or not to suppress its value or
88 # to keep it.
89 print('Suppressing pixels...')
90 G_mag_NMS = np.zeros([size[0], size[1]])
91 G_mag_NMS = np.copy(G_mag)
92 for r in range(0, size[0]):
93     for c in range(0, size[1]):
94         # Store the current pixel in a variable.
95         pixel = G_mag_pad[r + 1, c + 1]
96
97         # Grab the gradient angle.
98         grad_angle = np.radians(G_dir_round[r, c])
99
100         # Compute the row and columns offsets given by the direction matrix.
101         direction_y_offset = int(np.round(np.sin(grad_angle)))
102         direction_x_offset = int(np.round(np.cos(grad_angle)))
103
104         # Calculate which pixels are neighboring the current pixel based on the offsets.
105         G_north = G_mag_pad[r + 1 + direction_y_offset, c + 1 + direction_x_offset]
106         G_south = G_mag_pad[r + 1 - direction_y_offset, c + 1 - direction_x_offset]
107
108         # If the current pixel is greater than its neighbors, do not suppress it.
109         if pixel > G_north and pixel > G_south:
110             G_mag_NMS[r, c] = pixel
111         else:
112             G_mag_NMS[r, c] = 0
113
114 # Start the thresholding process.
115 print('-----')
116 print('---Thresholding---')
117 G_mag_NMS = (255 / np.max(G_mag_NMS)) * G_mag_NMS
118 G_threshold = np.where(G_mag_NMS > threshold, 255, 0)
119
120 # Show some of the intermediate steps.
121 print('-----')
122 print('---Displaying---')
123
124 # Show the gradient magnitude image.
125 G_mag_img = np.uint8(G_mag * (255 / np.max(G_mag)))
126 cv2.imshow('Gradient Magnitude', G_mag_img)
127 cv2.imwrite('gradient_magnitude.jpg', G_mag_img)
128 cv2.waitKey(0)
129
130 # Show the NMS gradient magnitude image.
131 G_mag_NMS_img = np.uint8(G_mag_NMS)
132 cv2.imshow('NMS Gradient Magnitude', G_mag_NMS_img)
133 cv2.imwrite('nms_gradient_magnitude.jpg', G_mag_NMS_img)
134 cv2.waitKey(0)
135
136 # Return the final edged image.
137 return np.uint8(G_threshold)
138
139
140 # Function to compute the convolution between an image and a given kernel.
141 def compute_convolution(im, kernel):
142     # Grab the sizes of the given image and kernel.
143     size_im = im.shape
144     size_k = kernel.shape
145
146     # Grab the offset generated by the kernel assuming the kernel is square with odd dimensionality.
147     o = int(np.floor(size_k[0] / 2))
148
149     # Pad the image based on the dimensions of the kernel.
150     im_pad = np.pad(im, o, mode='symmetric')
151
152     # Compute the convolution between the kernel and the padded image.
153     im_smooth = np.zeros([size_im[0], size_im[1]])
154     for r in range(o, size_im[0]):
155         for c in range(o, size_im[1]):
156             im_smooth[r, c] = np.multiply(kernel, im_pad[(r - o):(r + o + 1), (c - o):(c + o + 1)]).sum()
157
158     # Return the convolved image.
159     return im_smooth
160
161
162 # Function to round a given degree value to the nearest 45 degrees.
163 def round_degrees(degree):
164     # Correct the degree value first.

```

```

165     if degree < 0:
166         degree = degree + 360
167
168     # Provide a list of valid degree measurements.
169     valid_degrees = [0, 45, 90, 135, 180, 225, 270, 315, 360]
170
171     # Return the corrected degree value.
172     return valid_degrees[np.argmax(abs(valid_degrees - degree))]
173
174
175 # Load the geisel.jpg image.
176 file_path = r"HW3_geisel.jpg"
177 A = cv2.cvtColor(cv2.imread(file_path), cv2.COLOR_BGR2GRAY)
178 size = A.shape
179
180 # Compute the canny edge image.
181 A_edges = compute_canny_edge(A, 50)
182
183 # Show the original image.
184 cv2.imshow('Original Image', A)
185 cv2.imwrite('original_geisel.jpg', A)
186 cv2.waitKey(0)
187
188 # Resize the edged image.
189 cv2.imshow('Canny Edge Image', A_edges)
190 cv2.imwrite('canny_edge.jpg', A_edges)
191 cv2.waitKey(0)

```

Output:


Original Image Color:



Original Image Grayscale:



Gradient Image:



Gradient-NMS Image:





Image after Thresholding:



Here, a thresholding intensity value of **50** seemed sufficient to adequately display the obvious structure edges. It was difficult to strike a balance between removing noise produced by foliage while also preserving more obvious, thin edges such as the edge between the grass and the sidewalk. However, with this value, it is still clear that there exists some edge between the grass and concrete, and most of the noisier foliage has been removed from the image.

Problem 2:

Code:

```
194 #-----
195 #-----
196 # PROBLEM 1
197 #-----
198 cv2.destroyAllWindows()
199
200
201 # Define a function that will automatically return the log-shifted fourier transform of a
202 # given image.
203 def compute_fft(im):
204     # First, pad the image to a size of 512 x 512 pixels.
205     size = im.shape
206     im_pad = np.zeros([512, 512])
207     im_pad[0:size[0], 0:size[1]] = im
208
209     # Now, compute the fft of the padded image.
210     im_fft = np.fft.fft2(np.uint8(im_pad))
211
212     # Shift the fft.
213     im_fft_shift = np.fft.fftshift(im_fft)
214
215     # Take the log of the magnitude of the image.
216     im_fft_shift_log = np.log(np.abs(im_fft_shift))
217
218     # Return the image.
219     return [np.uint8(im_fft_shift_log), im_fft_shift]
220
221
222 # Function which accepts the parameters n, d0, and center coordinates u and v to generate the
223 # Butterworth notch filter mask of given size.
224 def generate_butterworth(size, n, D0, U, U_k, V, V_k):
225     # Generate an initial mask of all ones.
226     mask = np.ones([size[0], size[1]])
227     for u in range(0, size[0]):
228         for v in range(0, size[1]):
229             # Calculate the distances of the current point (u, v) to the specified centers of
230             # the Butterworth filter.
231             Dk_pos = np.sqrt((U[u, v] - U_k + 0.01)**2 + (V[u, v] - V_k + 0.01)**2)
232             Dk_neg = np.sqrt((U[u, v] + U_k + 0.01)**2 + (V[u, v] + V_k + 0.01)**2)
233
234             # Calculate the two terms in the product.
235             term_1 = 1 / (1 + (D0 / Dk_pos)**(2*n))
236             term_2 = 1 / (1 + (D0 / Dk_neg)**(2*n))
237
238             # Calculate the product for every center specified.
239             mask[u, v] = np.prod(term_1 * term_2)
240
241     # Return the calculated mask.
242     return mask
243
244
245 # Define the information necessary for both images.
246 file_path = ["Car.tif", "Street.png"]
247 u_k_dictionary = {0: [-85, -85, -85],
248                  1: [-166, 0]}
249
250 v_k_dictionary = {0: [-171, -85, 85, 171],
251                  1: [0, 166]}
252
253 n_dictionary = {0: 3,
254                1: 3}
255
256 D0_dictionary = {0: 15,
257                 1: 20}
258
259 # For both the Car and Street images, filter them with the Butterworth notch filter.
260 for n in range(0, 2):
261     print('Computing information for ' + file_path[n] + '...')
262
263     # Load the specified image.
264     A = cv2.imread(file_path[n])
265     A = A[:, :, 0]
266     image_size = A.shape
267
268     # Show the original image.
269     plt.figure(1)
270     plt.imshow(A, cmap='gray')
271     plt.colorbar()
272     plt.xlabel('x')
273     plt.ylabel('y')
274     plt.title('Original Image for ' + file_path[n])
275
276     # Compute the fourier transform.
277     print('Compute fft...')
278     [A_fft_log, A_fft] = compute_fft(A)
279     print('Done!')
280
281     # Show the fft of the image.
282     plt.figure(2)
283     plt.imshow(15 * A_fft_log, cmap='gray')
284     plt.colorbar()
285     plt.xlabel('u')
286     plt.ylabel('v')
287     plt.title('2D Log-Magnitude DFT Image for ' + file_path[n])
```



```

288
289 # Generate the u and v values for the Butterworth filter.
290 x_axis = np.linspace(-256,255,512)
291 y_axis = np.linspace(-256,255,512)
292 [u,v] = np.meshgrid(x_axis,y_axis)
293
294 # Generate the approximate (u, v) coordinates at which the impulses are located.
295 u_k = u_k_dictionary[n]
296 v_k = v_k_dictionary[n]
297
298 # Generate the butterworth mask for the image.
299 print('Generating mask...')
300 mask = generate_butterworth([512, 512], n_dictionary[n], D_0_dictionary[n], u, u_k, v, v_k)
301 print('Done!')
302
303 # Show the generated mask.
304 plt.figure(3)
305 plt.imshow(np.uint8(255 * mask), cmap='gray')
306 plt.colorbar()
307 plt.xlabel('u')
308 plt.ylabel('v')
309 plt.title('Butterworth Notch Filter Mask for ' + file_path[n])
310
311 # Apply the mask to the DFT of the image.
312 print('Filtering with mask...')
313 A_fft_masked = A_fft * mask
314
315 # Compute the filtered image.
316 A_filtered = np.real(np.fft.ifft2(np.fft.ifftshift(A_fft_masked)))
317 A_filtered = A_filtered[0:image_size[0], 0:image_size[1]]
318
319 # Threshold the image by 255 such that pixels over this value are reduced to 255.
320 # This prevents artifacts in the uint8 version of the image.
321 A_filtered = np.where(A_filtered > 255, 255, A_filtered)
322 print('Done!')
323
324 # Show the mask over the DFT.
325 plt.figure(4)
326 plt.imshow(np.uint8(15 * (A_fft_log * mask)), cmap='gray')
327 plt.colorbar()
328 plt.xlabel('u')
329 plt.ylabel('v')
330 plt.title('Butterworth Notch Filter Mask Overlay for ' + file_path[n])
331
332 # Show the filtered image.
333 plt.figure(5)
334 plt.imshow(np.uint8(A_filtered), cmap='gray')
335 plt.colorbar()
336 plt.xlabel('x')
337 plt.ylabel('y')
338 plt.title('Butterworth Filtered Image for ' + file_path[n])
339 plt.show()


```

Output:

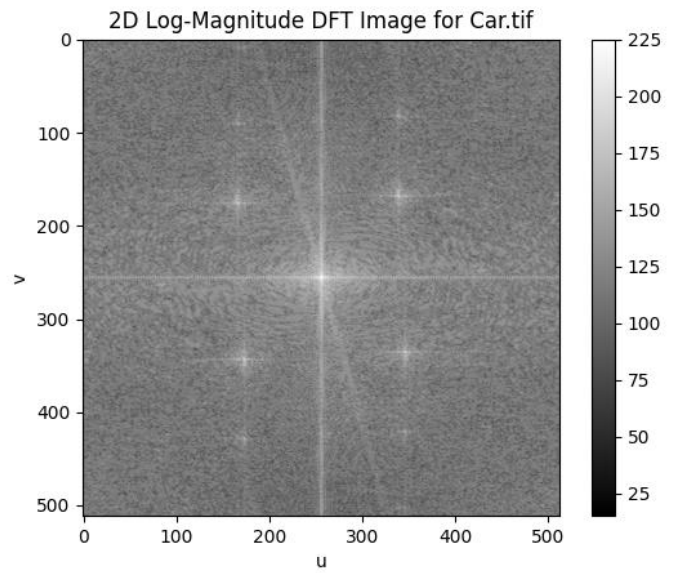
Part i:

Table of Parameters	
n	3
D_0	15
(u_1, v_1)	$(-85, -171)$
(u_2, v_2)	$(-85, -85)$
(u_3, v_3)	$(-85, 85)$
(u_4, v_4)	$(-85, 171)$

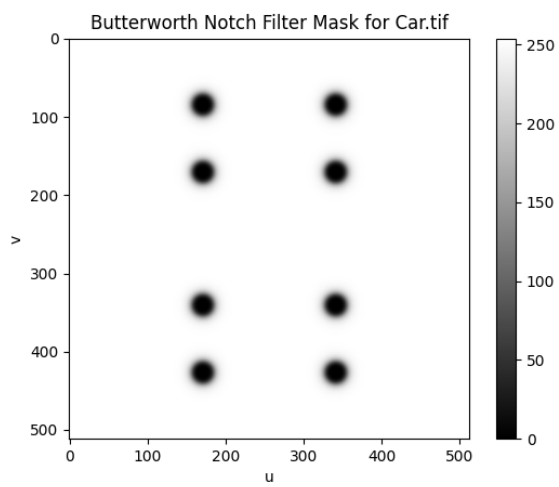
Original Image:



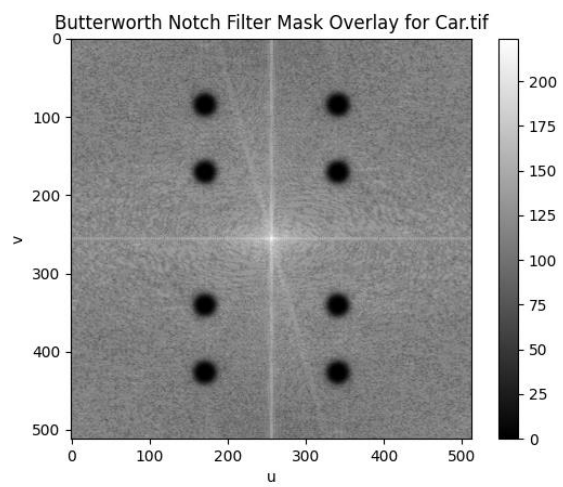
DFT Image:




Filter Mask



Filter Mask Overlayed on DFT



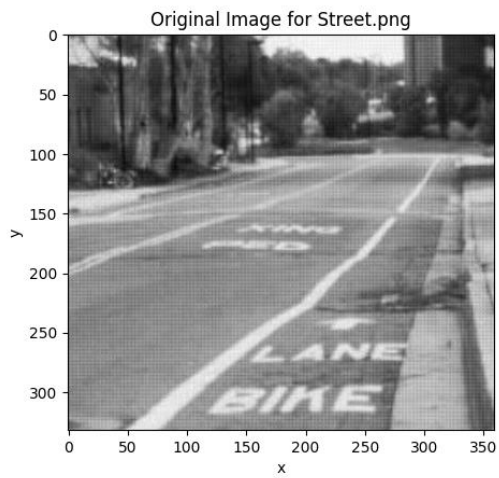
Filtered Image



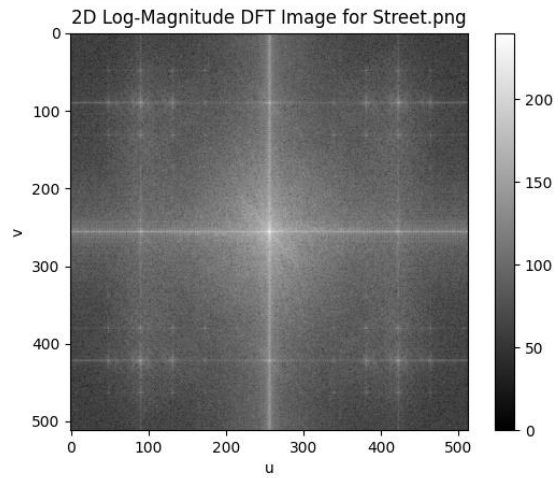
Part ii:

Table of Parameters	
n	3
D_0	20
(u_1, v_1)	$(-166, 0)$
(u_2, v_2)	$(0, 166)$


Original Image:




DFT Image:




Zoomed-in:



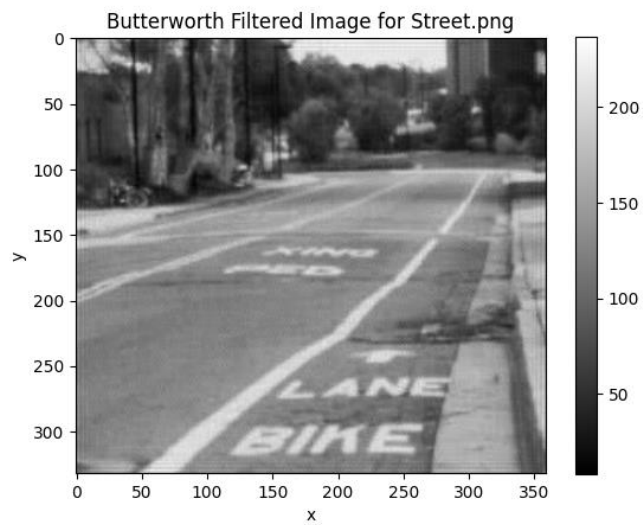
Filter Mask




Filter Mask Overlayed on DFT



Filtered Image



Zoomed-in:



Problem 3:

Code:

```
22 #-----
23 # Prep work.
24 #-----
25
26 # Define the transform from image data to tensor data.
27 transform = transforms.Compose([transforms.ToTensor(),
28                                transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
29
30 # Answer question ii.
31 print('Question ii\n-----')
32 print('Here, we are normalizing the images. However, we are not necessarily normalizing the images')
33 print('Indirectly. We are converting the images in the dataset to Tensors, then we are normalizing')
34 print('the Tensors to [-1, 1].')
35
36 # Define the batch size.
37 batch_size = 4
38
39 # Define the training set for the classifier.
40 trainset = torchvision.datasets.CIFAR10(root='./data',
41                                         train=True,
42                                         download=False,
43                                         transform=transform)
44
45 # Answer question i.
46 print('\nQuestion i\n-----')
47 print('There are ' + str(len(trainset)) + ' training images in the CIFAR10 dataset.')
48 print('Since we are using a batch size of ' + str(batch_size) + ', we have ' +
49       str(int(len(trainset) / batch_size)) + ' batches for training.')
50
51
52 # Load the training set for the classifier.
53 trainloader = torch.utils.data.DataLoader(trainset,
54                                           batch_size=batch_size,
55                                           shuffle=True,
56                                           num_workers=2)
57
58 # Define the testing set for the classifier.
59 testset = torchvision.datasets.CIFAR10(root='./data',
60                                         train=False,
61                                         download=False,
62                                         transform=transform)
63
64
65 # Print how many testing images there are.
66 print('There are ' + str(len(testset)) + ' testing images in the CIFAR10 dataset.')
67
68
69 # Load the testing set for the classifier.
70 testloader = torch.utils.data.DataLoader(testset,
71                                           batch_size=batch_size,
72                                           shuffle=True,
73                                           num_workers=2)
74
75 # Define the class identifiers for the classifier.
76 classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
77
78 #-----
79 # Showing some
80 # training images.
81 #-----
82
83 # Define an function to show a specified image.
84 def imshow(img):
85     # Unnormalize the image.
86     img = img / 2 + 0.5
87
88     # Convert the Tensor to a numpy array.
89     npimg = img.numpy()
90
91     # Show the image.
92     plt.imshow(np.transpose(npimg, (1, 2, 0)))
93     plt.show()
94
95 # Grab some random training images.
96 dataiter = iter(trainloader)
97 images, labels = dataiter.next()
98
99 # Show the random images.
100 imshow(torchvision.utils.make_grid(images))
101 print(' '.join('%5s' % classes[labels[j]] for j in range(batch_size)))
102
103 #-----
104 # Define a CNN
105 #-----
106
107 # Define a new class which will instantiate a CNN module.
108 class Net(nn.Module):
109     def __init__(self):
110         super().__init__()
111         self.conv1 = nn.Conv2d(3, 6, 5)
112         self.pool = nn.MaxPool2d(2, 2)
113         self.conv2 = nn.Conv2d(6, 16, 5)
114         self.fc1 = nn.Linear(16 * 5 * 5, 120)
115         self.fc2 = nn.Linear(120, 84)
```

```

116         self.fc3 = nn.Linear(84, 10)
117
118     def forward(self, x):
119         x = self.pool(F.relu(self.conv1(x)))
120         x = self.pool(F.relu(self.conv2(x)))
121         x = torch.flatten(x, 1)
122         x = F.relu(self.fc1(x))
123         x = F.relu(self.fc2(x))
124         x = self.fc3(x)
125         return x
126
127 # Define a new CNN from our specified class.
128 net = Net()
129 print('CNN creation successful!')
130
131 #-----
132 # Define a loss
133 # function and
134 # optimizer.
135 #-----
136
137 # Define our loss criterion.
138 criterion = nn.CrossEntropyLoss()
139 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
140 print('Criterion established!')
141
142 #-----
143 # Train the
144 # network.
145 #-----
146
147 print('Starting to train...')
148 loss_data = []
149 for epoch in range(2):
150     # Instantiate our current loss.
151     running_loss = 0.0
152     for i, data in enumerate(trainloader, 0):
153         # Get the inputs where the data is a tuple of input images and labels.
154         inputs, labels = data
155
156         # Zero the parameter gradients.
157         optimizer.zero_grad()
158
159         # Forward + backward + optimize
160         outputs = net(inputs)
161         loss = criterion(outputs, labels)
162         loss.backward()
163         optimizer.step()
164
165         # Print the current statistics for the classification.
166         running_loss += loss.item()
167         if i % 2000 == 1999:
168             loss_data.append(running_loss / 2000)
169             print('%d, %5d] loss: %.3f' % (epoch + 1, i + 1, running_loss / 2000))
170             running_loss = 0.0
171
172 print('Finished Training!')
173
174 #-----
175 # Plot the data
176 # loss.
177 #-----
178
179 # Question iii.
180 batch_count = [2000, 4000, 6000, 8000, 10000, 12000, 14000, 16000, 18000, 20000, 22000, 24000]
181 batch_labels = [2000, 4000, 6000, 8000, 10000, 12000, 14000, 16000, 18000, 20000, 22000, 24000]
182 batch = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12]
183 epoch_labels = [0, 2]
184
185 figure, axes = plt.subplots()
186 axes.set_xticks(batch)
187 axes.set_xticklabels(batch_labels)
188 axes.set_xlabel('Batch Count')
189 axes.set_ylabel('Data Loss')
190
191 epoch_axis = axes.twinx()
192 epoch_axis.set_xlim(0, 4)
193 epoch_axis.set_xticks(epoch_labels)
194 epoch_axis.set_xticklabels([1, 2])
195 epoch_axis.set_xlabel('Epochs')
196
197 axes.plot(batch, loss_data, '-bo')
198 plt.title('Training Loss for the Network')
199 plt.show()
200 plt.rcParams['figure.figsize'] = [15, 8]
201
202 #-----
203 # Save the
204 # trained model.
205 #-----
206 PATH = './cifar_net.pth'
207 torch.save(net.state_dict(), PATH)

```

```

208
209 #-----
210 # Test the
211 # network on the
212 # loaded test
213 # data.
214 #-----
215
216 # Instantiate a iterator for the loaded test data.
217 dataiter = iter(testloader)
218 images, labels = dataiter.next()
219
220 # Print some of the test images.
221 imshow(torchvision.utils.make_grid(images))
222 print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))
223
224 # Load a CNN with the trained model.
225 net = Net()
226 net.load_state_dict(torch.load(PATH))
227
228 # Generate the outputs.
229 outputs = net(images)
230 _, predicted = torch.max(outputs, 1)
231 print('Predicted: ', ' '.join('%5s' % classes[predicted[j]] for j in range(4)))
232
233 # Calculate the accuracy of the network.
234 correct = 0
235 total = 0
236 print('Calculating the accuracy of the network...')
237 with torch.no_grad():
238     for data in testloader:
239         images, labels = data
240         outputs = net(images)
241         _, predicted = torch.max(outputs.data, 1)
242         total += labels.size(0)
243         correct += (predicted == labels).sum().item()
244
245 print('Accuracy of the network on the 10000 test images: %d %%' % (100 * correct / total))
246
247 # prepare to count predictions for each class
248 correct_pred = {classname: 0 for classname in classes}
249 total_pred = {classname: 0 for classname in classes}
250
251 print('Calculating accurate predictions per label...')
252 with torch.no_grad():
253     for data in testloader:
254         images, labels = data
255         outputs = net(images)
256         _, predictions = torch.max(outputs, 1)
257         # collect the correct predictions for each class
258         for label, prediction in zip(labels, predictions):
259             if label == prediction:
260                 correct_pred[classes[label]] += 1
261                 total_pred[classes[label]] += 1
262
263 print('Printing accuracy for each label...')
264 for classname, correct_count in correct_pred.items():
265     accuracy = 100 * float(correct_count) / total_pred[classname]
266     print("Accuracy for class {:5s} is: {:.1f} %".format(classname, accuracy))
267
268
269 # Instantiate a new loader which will only load one image at a time.
270 single_loader = torch.utils.data.DataLoader(trainset,
271                                             batch_size=1,
272                                             shuffle=True,
273                                             num_workers=2)
274
275 # Define a imshow method for a single layer in the Conv2 output.
276 def imshow_layer(img):
277     # Unnormalize the image.
278     img = img / 2 + 0.5
279
280     # Convert the Tensor to a numpy array.
281     npimg = img.numpy()
282
283     # Show the image.
284     plt.imshow(npimg, cmap='gray')
285     plt.show()
286
287 # Since the data is shuffled, grab the first random image.
288 for index, data in enumerate(single_loader, 0):
289     temp = data[index]
290     imshow(torchvision.utils.make_grid(temp))
291     out = F.relu(net.conv1(temp))
292     break
293
294 # For every layer in the Tensor, plot it as a grayscale image.
295 [blank, layers, r, c] = out.shape
296 for i in range(0, layers):
297     print('Feature Map Layer ' + str(i + 1))
298     imshow_layer(out[0, i, :, :].detach())

```


Output:

Question i:

Question i

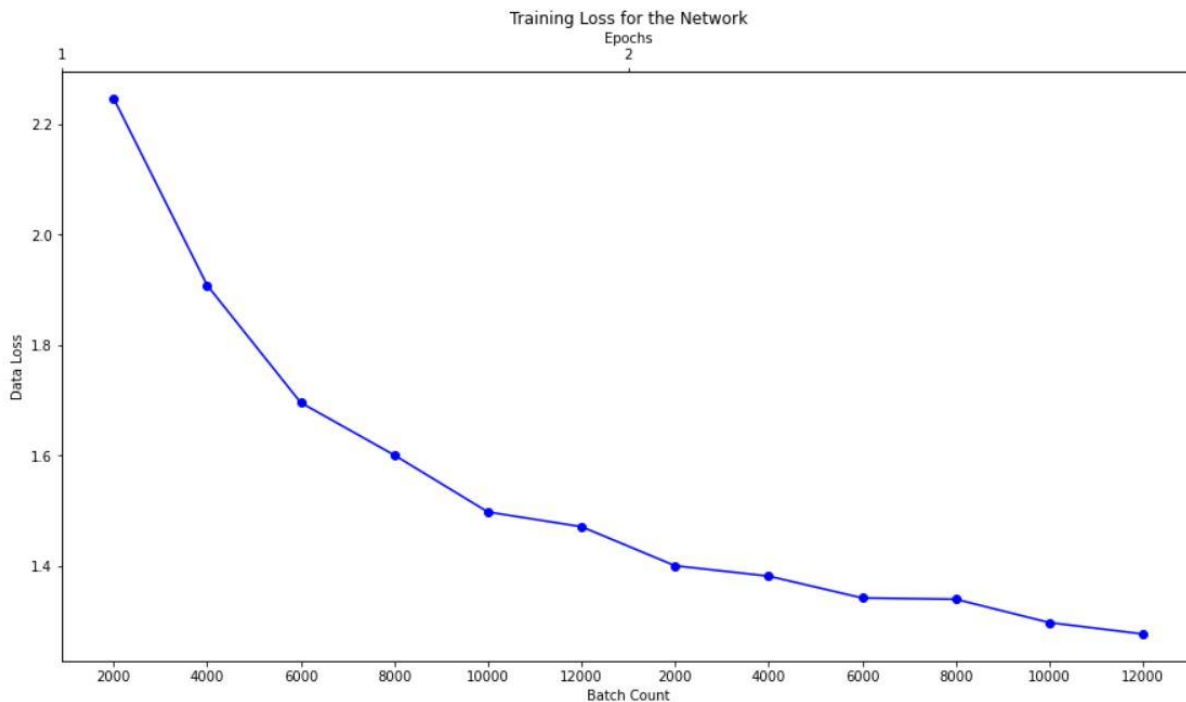
There are 50000 training images in the CIFAR10 dataset.
Since we are using a batch size of 4, we have 12500 batches for training.
There are 10000 testing images in the CIFAR10 dataset.

Question ii:


Question ii

Here, we are normalizing the images. However, we are not necessarily normalizing the images directly. We are converting the images in the dataset to Tensors, then we are normalizing the Tensors to $[-1, 1]$.


Question iii:



Question iv:



GroundTruth: deer deer car truck
Predicted: deer bird car truck




GroundTruth: horse cat truck horse
Predicted: cat horse frog horse

Here, we can see that the predicted category for the given image does not always match with the ground truth (i.e., the label assigned to the image before passed through the network). For example, in the second row, the first image was predicted to be a cat, when the image is actually of a horse. It might be likely that this discrepancy is a result of someone appearing to be riding the horse.


Question v:

Original Image:




Feature Maps Produced by the First Convolution:


Feature Map Layer 1




Feature Map Layer 2




Feature Map Layer 3



Feature Map Layer 4



Feature Map Layer 5



Feature Map Layer 6

