

Sean Carda
PID: A59009786
ECE 253 – Image Processing
December 8, 2021
Homework 4

Academic Integrity Policy: Integrity of scholarship is essential for an academic community. The University expects that both faculty and students will honor this principle and in so doing protect the validity of University intellectual work. For students, this means that all academic work will be done by the individual to whom it is assigned, without unauthorized aid of any kind.

By including this in my report, I agree to abide by the Academic Integrity Policy mentioned above.

Problem 1:

Code:

```
1 #-----
2 # Sean Carda
3 # ECE 253 - Image Processing
4 # Dr. Mohan Trivedi
5 # Homework 3
6 #-----
7
8 #-----
9 # Global Imports.
10 import cv2
11 import numpy as np
12 import scipy.ndimage as ndi
13 import matplotlib.pyplot as plt
14 #-----
15 # Suppress scientific notation representation for numpy arrays.
16 np.set_printoptions(suppress=True)
17
18 #-----
19
20 # Cross Correlation Filter.
21
22 # First, read in the first birds image and the template we will use to identify the features
23 # of interest in the image.
24 print('Loading images...')
25 birds = cv2.cvtColor(cv2.imread('birds1.jpeg'), cv2.COLOR_BGR2GRAY)
26 birds_2 = cv2.cvtColor(cv2.imread('birds2.jpeg'), cv2.COLOR_BGR2GRAY)
27 template = cv2.cvtColor(cv2.imread('template.jpeg'), cv2.COLOR_BGR2GRAY)
28 print('Done!')
29
30 # Define a method for showing images which will reduce the lines of code needed
31 # to show images.
32 def imshow(im, title, fig_num, xlabel, ylabel):
33     plt.figure(fig_num)
34     plt.imshow(im, cmap='gray')
35     plt.colorbar()
36     plt.title(title)
37     plt.xlabel(xlabel)
38     plt.ylabel(ylabel)
39
40 # Show the original image and template.
41 imshow(birds, 'Original Birds 1 Image', 1, 'Columns', 'Rows')
42 imshow(birds_2, 'Original Birds 2 Image', 2, 'Columns', 'Rows')
43 imshow(template, 'Template for Matching', 3, 'Columns', 'Rows')
44
45 # Compute convolution between the birds image and the template.
46 print('Cross correlating...')
47 matching = ndi.convolve(np.float64(birds), np.float64(template))
```

```


48 imshow(matching, 'Birds 1 Image Convolved with Template', 4, 'Columns', 'Rows')
49 print('Done!')
50 plt.show()
51
52 #-----
53 # Normalized Cross Correlation Filter.
54
55 # Define a function that accepts an image and a template and computes the normalized
56 # cross-correlated image.
57 def normalized_cc(im, template):
58     # Grab the sizes of the image and template.
59     im_size = im.shape
60     t_size = template.shape
61
62     # Pad the image with the size of the template.
63     im_pad = np.pad(im, (int(t_size[0] / 2) + 1, int(t_size[1] / 2)), mode='symmetric')
64
65     # Instantiate an output image.
66     output = np.zeros(im_size)
67
68     # Calculate the mean of the template and the term of the normalization contributed by the
69     # template.
70     t_avg = np.mean(template)
71     t_offset = template - t_avg
72     norm_coeff_t = np.sum((template - t_avg)**2)
73     for r in range(0, im_size[0]):
74         for c in range(0, im_size[1]):
75             # Grab the image windowed by the template.
76             im_window = im_pad[r:r+t_size[0], c:c+t_size[1]]
77
78             # Compute the average of the image in the current window.
79             im_avg = np.mean(im_window)
80
81             # Using the calculated means, compute the cross correlation of the image in the current
82             # window.
83             im_offset = (im_window - im_avg)
84             cross_corr = np.sum(im_offset * t_offset)
85
86             # Compute the normalized coefficients of the image and the template at the current
87             # window.
88             norm_coeff_im = np.sum(im_offset**2)
89             norm_coeff = norm_coeff_t * norm_coeff_im
90
91             # Compute the normalized cross-correlation.
92             norm_cc = cross_corr / np.sqrt(norm_coeff)
93             output[r, c] = norm_cc
94
95     # Return.
96     print('Done!')
97     return output
98
99
100 # Accepts an image and a point in the image and places a box of given dimension around that
101 # point in the image.
102 def identify(im, point, dim):
103     # Define an output image which will be boxed.
104     size = im.shape
105     im_boxed = np.zeros([size[0], size[1], 3])
106     im_boxed[:, :, 0] = np.copy(im)
107     im_boxed[:, :, 1] = np.copy(im)
108     im_boxed[:, :, 2] = np.copy(im)
109
110     # Calculate the height and width of the box.
111     width = int(dim[1] / 2)
112     height = int(dim[0] / 2)
113
114     # Create points.
115     point_1 = (point[1] - width, point[0] - height)
116     point_2 = (point[1] + width, point[0] + height)
117
118     # Draw the rectangle.
119     cv2.rectangle(im_boxed, point_1, point_2, color=(255, 0, 0), thickness=3)
120
121     # Return.
122     return np.uint8(im_boxed)
123
124 # Compute the normalized cross correlation on the first birds image.
125 print('Computing normalized cross correlation for birds 1...')
126 birds_norm_cc = normalized_cc(np.float64(birds), np.float64(template))
127 imshow(birds_norm_cc, 'Birds Normalized Cross Correlation', 1, 'Columns', 'Rows')
128
129 # Place a box around the point of strongest intensity.
130 point = np.unravel_index(np.argmax(birds_norm_cc), birds_norm_cc.shape)
131 print(point)
132 matched_image = identify(birds, point, template.shape)
133 imshow(matched_image, 'Birds 1 with Identified Template Match', 2, 'Columns', 'Rows')
134
135 # Compute the normalized cross correlation on the second birds image.
136 print('Computing normalized cross correlation for birds 2...')
137 birds_norm_cc = normalized_cc(np.float64(birds_2), np.float64(template))
138 imshow(birds_norm_cc, 'Birds 2 Normalized Cross Correlation', 3, 'Columns', 'Rows')
139
140 # Place a box around the point of strongest intensity.
141 point = np.unravel_index(np.argmax(birds_norm_cc), birds_norm_cc.shape)
142 print(point)
143 matched_image = identify(birds_2, point, template.shape)
144 imshow(matched_image, 'Birds 2 with Identified Template Match', 4, 'Columns', 'Rows')

```

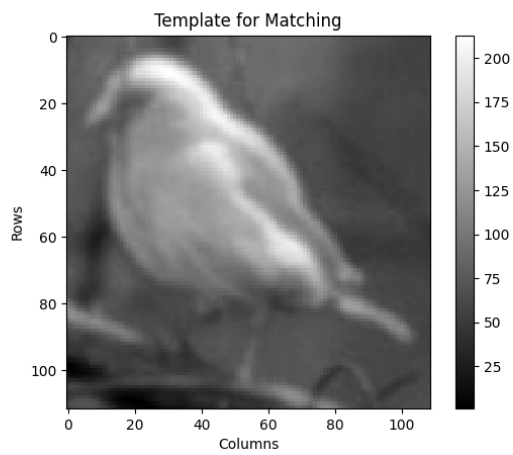
Output:

Cross Correlation:

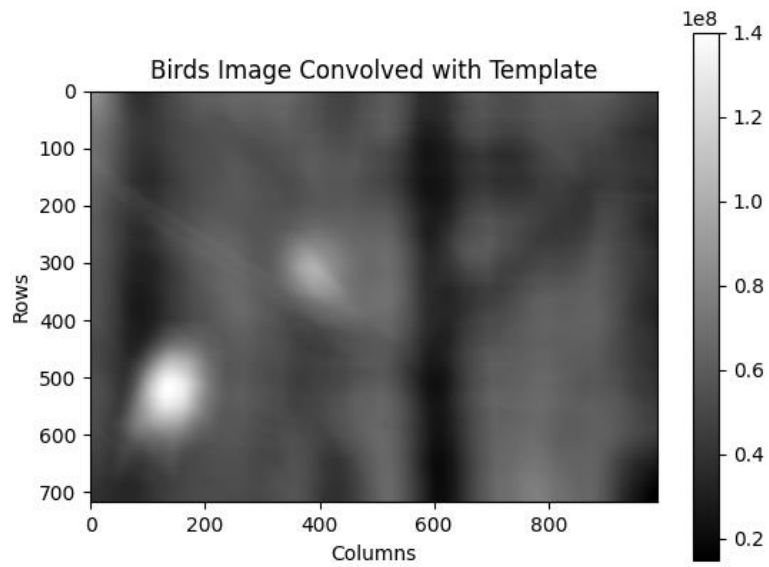
Original Birds 1 Image:



Template Image:

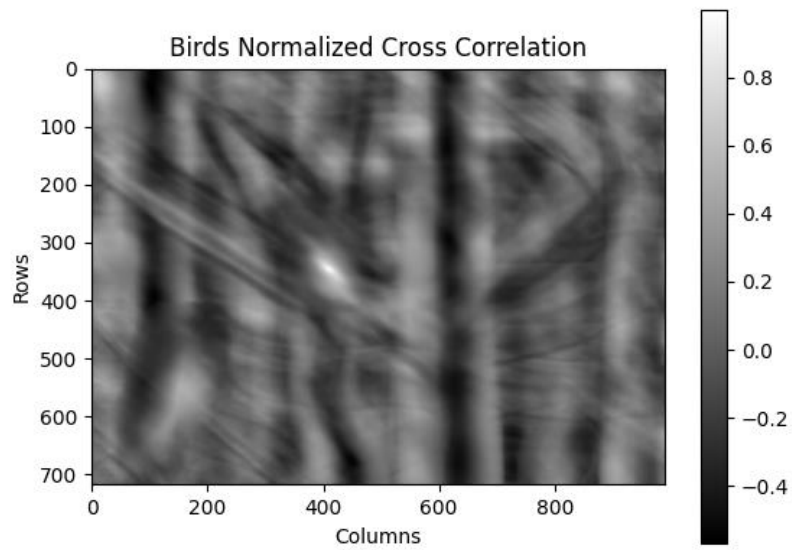


Cross Correlation of Birds 1 with Template:




Normalized Cross Correlation:

Birds 1:




Birds 1 Matched:




Here, it is evident that the template has found a match in the Birds 1 image. This is simple due to the fact that the bird's orientation and lightning in the image strongly matches with the bird in the template.


Original Birds 2 Image:



Birds 2 NCC:



Birds 2 Matched:



Here, it is clear that none of the birds have been identified in the image. This is due to the fact that the lighting and orientation of all the birds in the image do not match up with any components of the template image. It's possible that if one of the birds were facing the other direction, we would have a stronger correlation between the image and the template.

Problem 2

Code:

```
1 #-----
2 # Sean Carda
3 # ECE 253 - Image Processing
4 # Dr. Mohan Trivedi
5 # Homework 3
6 #-----
7
8 #-----
9 # Global Imports.
10 import cv2
11 import numpy as np
12 import scipy.ndimage as ndi
13 import matplotlib.pyplot as plt
14 #-----
15 # Suppress scientific notation representation for numpy arrays.
16 np.set_printoptions(suppress=True)
17
18 #-----
19 # HOUGH TRANSFORM ON TEST IMAGE.
20
21 # Method which accepts an image and computes the Hough transform.
22 def HT(im):
23     # Maximum angle for the transform.
24     t_max = 90
25
26     # Generate all possible theta values.
27     theta = np.int64(np.round(np.linspace(t_max, t_max, 2 * t_max + 1)))
28
29     # Calculate the maximum distance between opposite corners of the image.
30     size = im.shape
31     D = np.int64(np.floor(np.sqrt((size[0]**2) + (size[1]**2))))
32     rho = np.int64(np.round(np.linspace(-D, D, 2 * D + 1)))
33
34     # Instantiate an output image.
35     output = np.zeros([len(rho), len(theta)])
36
37     for r in range(0, size[0]):
38         for c in range(0, size[1]):
39             if im[r, c] >= 1:
40                 for t in theta:
41                     p = (im[r, c] * c * np.cos(np.deg2rad(t))) + (im[r, c] * r * np.sin(np.deg2rad(t)))
42                     p = np.int64(p + np.max(rho))
43                     output[p, t + t_max] += 1
44
45             if (r % 50) == 0:
46                 print('Progress: ' + str(round(100 * r / size[0], 2)) + '%...')
47
48     # Return the transform and the established theta and rho values.
49     return [output, rho, theta]
50
51 # Method which accepts an image and a list of rho and theta parameters and draws the corresponding
52 # x-y lines on the image. Parameters = [list of theta, list of rho]
53 def draw_lines(params, im, width, round):
54     # Create a new image which will overlay the lines on top of the image.
55     size = im.shape
56     output = np.zeros([size[0], size[1], 3])
57     output[:, :, 0] = im
58     output[:, :, 1] = im
59     output[:, :, 2] = im
60
61     # For every parameter combination, draw the corresponding line.
62     for combo in params:
63         # Grab theta and rho
64         theta = combo[0]
65         rho = combo[1]
66
67         # Compute the new line parameters.
68         if theta == 0:
69             theta = 0.01
70         elif theta == 90:
71             theta = 89.99
72         elif theta == -90:
73             theta = -89.99
74
75         b = rho / (np.sin(np.deg2rad(theta)))
76         if round is True:
77             m = np.round(-1 / (np.tan(np.deg2rad(theta))))
78         else:
79             m = -1 / (np.tan(np.deg2rad(theta)))
80
81         # Compute the two points needed to generate the line
82         point_1 = ((-999, int(np.round((-999 * m) + b))))
83         point_2 = ((999, int(np.round((999 * m) + b))))
84
85         # Draw the line onto the image.
86         cv2.line(output, point_1, point_2, (255, 0, 0), width, LineType=cv2.LINE_AA)
87
88     # Return the image.
89     return output
90
91 # Method which accepts a thresholded transform and the respective rho and theta values and
92 # returns the theta-row indices of the transformed image.
93 def get_params(m, rho, theta):
94     # Return the indices at which the transform is nonzero.
95     list = np.transpose(np.nonzero(m))
```

```

96
97     # Augment the list by the theta values by the maximum theta value.
98     list[:, 1] = list[:, 1] - np.max(theta)
99     list[:, 0] = list[:, 0] - np.max(rho)
100
101     # Return the flipped array so that theta is an index [-, 0] and rho is at [-, 1]
102     list_flip = np.flip(list, 1)
103     return list_flip
104
105
106 #-----
107 # HOUGH TRANSFORM ON TEST IMAGE.
108 # First, generate an 11 x 11 image with 5 1's located in the center and in the four corners.
109 test_im = np.zeros([11, 11])
110 size = test_im.shape
111 points = [[0, 0], [0, size[1] - 1], [int(size[0] / 2), int(size[1] / 2)], [size[0] - 1, 0], [size[0] - 1, size[1] - 1]]
112 for point in points:
113     test_im[point[0], point[1]] = 1
114
115 # Calculate the Hough transform of the test image.
116 [transform, rho, theta] = Ht(test_im)
117
118 # Threshold the transformed image and look for intersections greater than 2.
119 transform_thresh = np.where(transform > 2, transform, 0)
120
121 # Draw the lines specified by the transform.
122 params = get_params(transform_thresh, rho, theta)
123 print(params)
124 test_lines = draw_lines(params, test_im, 1, True)
125
126 plt.figure(1)
127 plt.imshow(test_im, cmap='gray')
128 plt.title('Original Test Image')
129 plt.xlabel('y')
130 plt.ylabel('x')
131 plt.colorbar()
132
133 plt.figure(2)
134 plt.imshow(transform, extent=[theta[0], theta[len(theta) - 1], rho[len(rho) - 1], rho[0]], cmap='gray')
135 plt.title('Hough Transform on Test Image')
136 plt.xlabel('theta (degrees)')
137 plt.ylabel('p')
138 plt.colorbar()
139
140 plt.figure(3)
141 plt.plot(3, 5, 1, 9, color='white', linewidth=1)
142 plt.imshow(test_lines, cmap='gray')
143 plt.title('Test Image with Lines')
144 plt.xlabel('y')
145 plt.ylabel('x')
146 plt.colorbar()
147 plt.show()
148
149 #-----
150 # HOUGH TRANSFORM ON LANE IMAGE III
151
152 # Read in the lane image.
153 lane = cv2.imread('lane.png')
154 thresh = 240
155 binary_lane = cv2.Canny(lane[:, :, 0], 175, thresh, apertureSize=3, L2gradient=True) / 255
156
157
158 plt.figure(4)
159 plt.imshow(lane, cmap='gray')
160 plt.title('Original Lane Image')
161 plt.xlabel('y')
162 plt.ylabel('x')
163 plt.colorbar()
164
165 plt.figure(5)
166 plt.imshow(binary_lane, cmap='gray')
167 plt.title('Original Lane Image')
168 plt.xlabel('y')
169 plt.ylabel('x')
170 plt.colorbar()
171 plt.show()
172
173 # Compute the Hough transform of the edged image.
174 [transform, rho, theta] = Ht(np.uint8(binary_lane))
175
176 # Threshold the transformed image and look for intersections greater than 2.
177 transform_thresh = np.where(transform > (0.75 * np.max(transform)), transform, 0)
178
179 # Draw the lines specified by the transform.
180 params = get_params(transform_thresh, rho, theta)
181 print('Parameters for analysis:')
182 print(params)
183 print('Drawing Lines...')
184 lane_lines = draw_lines(params, lane[:, :, 0], 3, False)
185
186 plt.figure(4)
187 plt.imshow(lane, cmap='gray')
188 plt.title('Original Lane Image')
189 plt.xlabel('y')
190 plt.ylabel('x')
191 plt.colorbar()

```



```


192 plt.figure(5)
193 plt.imshow(binary_lane, cmap='gray')
194 plt.title('Edged Lane Image')
195 plt.xlabel('y')
196 plt.ylabel('x')
197 plt.colorbar()
198
199
200
201 plt.figure(6, figsize=(10, 20))
202 plt.imshow(transform, extent=[theta[0], theta[len(theta) - 1], rho[len(rho) - 1], rho[0]], aspect='auto', cmap='gray')
203 plt.title('Hough Transform on Lane Image')
204 plt.xlabel('theta (degrees)')
205 plt.ylabel('p')
206 plt.colorbar()
207
208 plt.figure(8)
209 plt.imshow(np.uint8(lane_lines), cmap='gray')
210 plt.title('Lane Image with Lines')
211 plt.xlabel('y')
212 plt.ylabel('x')
213 plt.colorbar()
214 plt.show()
215
216
217 #-----
218 # HOUGH TRANSFORM ON LANE IMAGE IV
219
220 # Set a theta range for which to preserve lane edges.
221 theta_n_low = -57 + 90
222 theta_n_high = -50 + 90
223 theta_p_low = 50 + 90
224 theta_p_high = 57 + 90
225
226 # Threshold the thresholded image again, but only preserving the ranges of theta specified.
227 t_size = transform_thresh.shape
228 theta_thresh = np.zeros(transform_thresh.shape)
229 for i in range(t_size[1]):
230     if (i > theta_n_low and i < theta_n_high) or (i > theta_p_low and i < theta_p_high):
231         theta_thresh[:, i] = transform_thresh[:, i]
232
233
234 # Draw the lines specified by the transform.
235 params = get_params(theta_thresh, rho, theta)
236 print('Parameters for analysis:')
237 print(params)
238 print('Drawing Lines...')
239 new_lane_lines = draw_lines(params, lane[:, :, 0], 3, False)
240
241 # Show the new lined image.
242 plt.figure(1)
243 plt.imshow(np.uint8(new_lane_lines), cmap='gray')
244 plt.title('New Lane Image with Lines')
245 plt.xlabel('y')
246 plt.ylabel('x')
247 plt.colorbar()
248 plt.show()

```

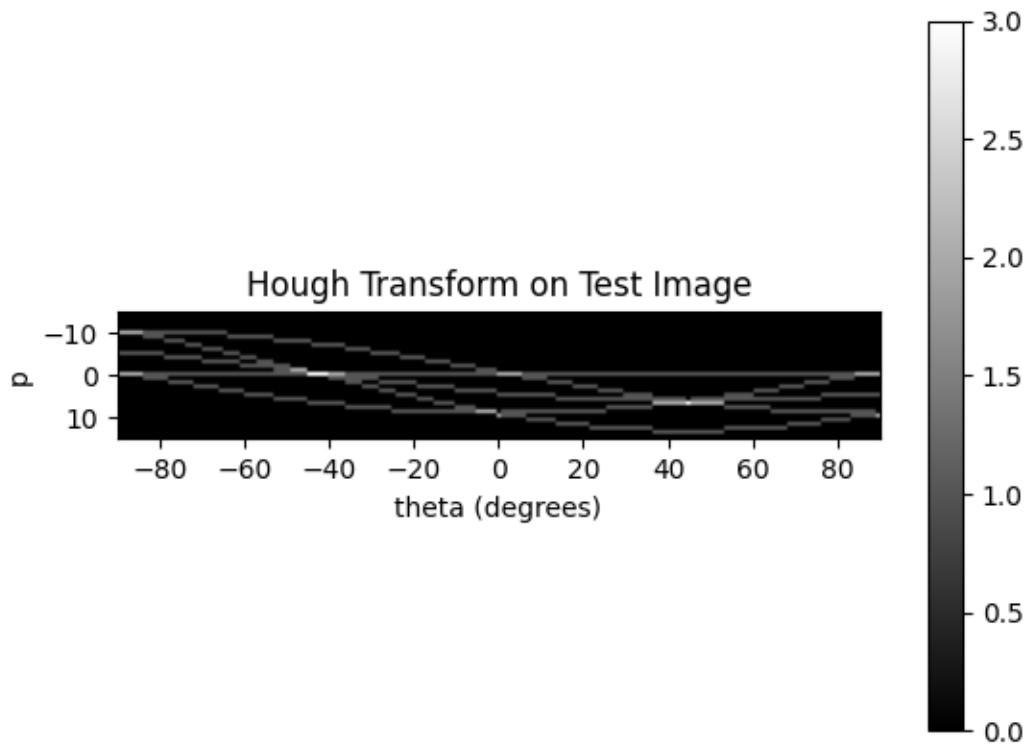
Output:

Test Image:

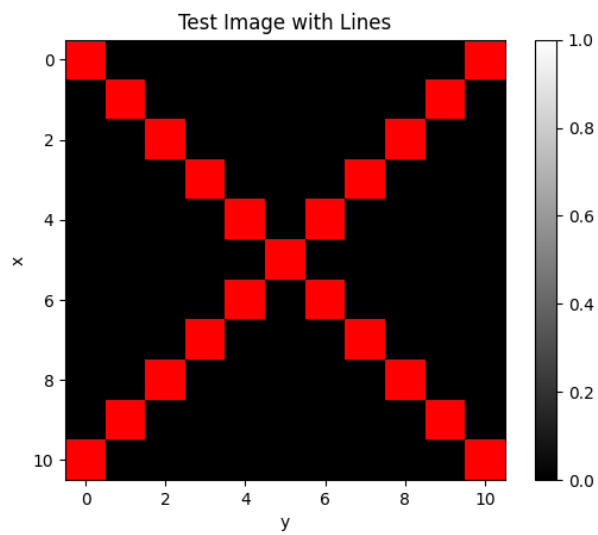
Original Image:




Hough Transform of Test Image:



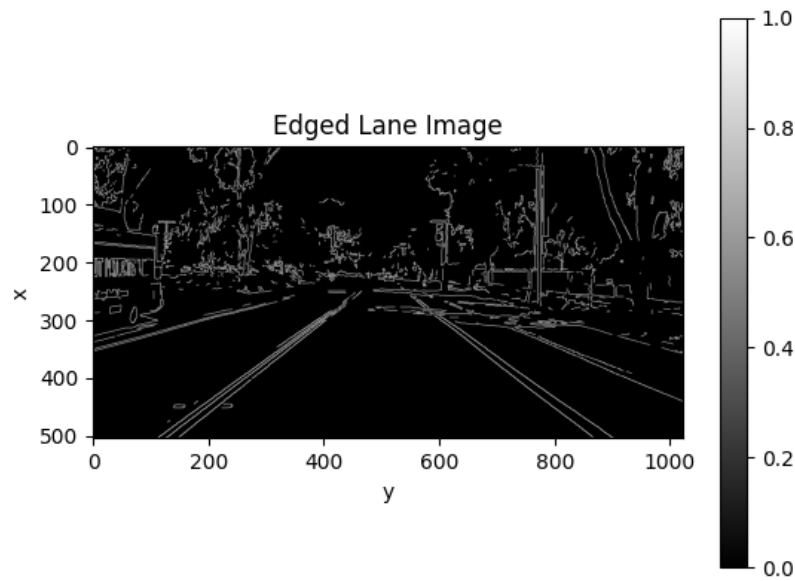
Lines Overlayed on Test Image:



Lanes Image:
Original Image:



Binary Edge Image:



Hough Transform of Edge Image:

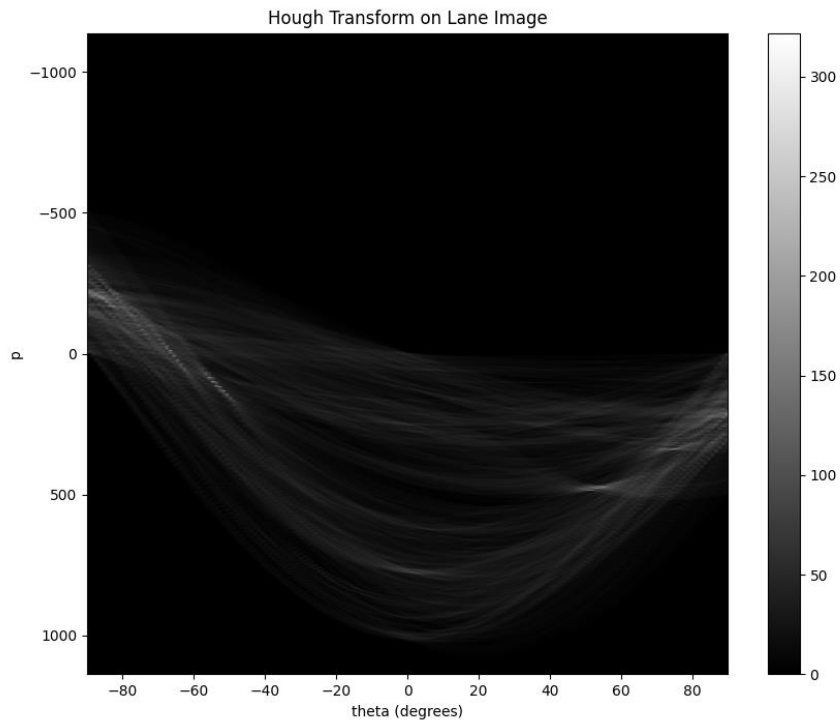



Image with Lines at 75% Threshold of HT:




Threshold values:

$$-57^\circ \leq \theta \leq -50^\circ$$

And

$$50^\circ \leq \theta \leq 57^\circ$$

Image with Lines at the Theta Threshold Above:



Problem 3:

Code:


```
1 #-----
2 # Sean Carda
3 # ECE 253 - Image Processing
4 # Dr. Mohan Trivedi
5 # Homework 3
6 #-----
7
8 #-----
9 # Global Imports.
10 import cv2
11 import numpy as np
12 import scipy.ndimage as ndi
13 import matplotlib.pyplot as plt
14 #-----
15 # Suppress scientific notation representation for numpy arrays.
16 np.set_printoptions(suppress=True)
17
18
19 # Method to create a dataset on which to cluster our data.
20 def createDataset(im):
21     # Grab the size of the image.
22     [N, M, D] = im.shape
23
24     # Separate the image into it's layered components.
25     red = im[:, :, 2]
26     green = im[:, :, 1]
27     blue = im[:, :, 0]
28
29     # Reshape the vectors to N*M x 1.
30     new_shape = (N * M, 1)
31     red_v = np.reshape(red, new_shape)
32     green_v = np.reshape(green, new_shape)
33     blue_v = np.reshape(blue, new_shape)
34
35     # Concatenate the new vectors into a dataset matrix.
36     features = np.zeros([N * M, 3])
37     features[:, 0] = blue_v[:, 0]
38     features[:, 1] = green_v[:, 0]
39     features[:, 2] = red_v[:, 0]
40
41     # Return.
42     return features
43
44
45 # Method to perform K-means segmentation on a given dataset.
46 def KMeansCluster(features, centers):
47     # Grab the length of the features and the length of the center matrices.
48     f_size = features.shape
49     q = f_size[0]
50     m_size = centers.shape
51     k = m_size[0]
52
53     # Instantiate a new cluster matrix which will store the pixels matching the center.
54     clusters = {}
55     difference = {}
56     for i in range(k):
57         clusters[i] = None
58         difference[i] = None
59
60     # Instantiate a set of new k-mean centers.
61     new_centers = centers
62
63     # Perform the k-means clustering by looping through the dataset and comparing
64     # pixels with k-mean centers.
65     iterations = 70
66     for iter in range(iterations):
67         clusters = {}
68         difference = {}
69         for j in range(k):
70             clusters[j] = None
71             difference[j] = None
72
73         # Compute the difference between the features and the centers.
74         for j in range(k):
75             difference[j] = (np.linalg.norm(features - new_centers[j, :], axis=1))**2
76
77         # Generate a matrix from the calculated norms.
78         diff_mat = np.zeros([q, k])
79         for j in range(k):
80             diff_mat[:, j] = difference[j]
81
82         # Compute the argmin of the rows of the difference matrix.
83         indices = np.argmin(diff_mat, axis=1)
84
85         # For the calculated indices, generate an array of the indices for a specified minimum
86         # index. Then, using those indices, find the features in those locations.
87         for j in range(k):
88             index_list = np.where(indices == j)
89             clusters[j] = np.array(features[index_list, :])
90             if clusters[j].size != 0:
91                 new_centers[j, :] = np.mean(clusters[j], axis=0)
92             else:
93                 new_centers[j, :] = np.random.randint(0, 255, size=(1, 3))
94
95     print(new_centers)
```

```


96
97     # Return the new calculated centers and clusters.
98     return [clusters, new_centers, indices]
99
100
101 # Method to map values in the old array to values calculated by k-means clustering.
102 def mapValues(im, idx, centers):
103     # Grab the various dimensions of the given objects.
104     [N, M, D] = im.shape
105     k_size = centers.shape
106     k = k_size[0]
107
108     # Replace the values of the old image with the k-mean centers. For every k,
109     # find the values in the idx array that match the current minimum index. Find the
110     # center that matches this index and replace the rows of the output with this value.
111     output = np.zeros([N * M, 3])
112     for i in range(k):
113         index_list = np.argwhere(idx == i)
114         if index_list.size != 0:
115             output[index_list, :] = centers[i, :]
116
117     # Extract color components.
118     red = np.reshape(output[:, 0], (N, M))
119     green = np.reshape(output[:, 1], (N, M))
120     blue = np.reshape(output[:, 2], (N, M))
121
122     # Generate the image from the extracted color components.
123     new_im = np.zeros(im.shape)
124     new_im[:, :, 0] = blue
125     new_im[:, :, 1] = green
126     new_im[:, :, 2] = red
127
128     # Return the image.
129     return new_im
130
131
132 # Method to generate a k x 3 set of random k-mean centers.
133 def generateMeans(k):
134     # Instantiate a random set of k means (centers).
135     return np.random.randint(20, 200, size=(k, 3))
136
137
138 # Read in an image.
139 file_path = r"white-tower.png"
140 #A = cv2.cvtColor(cv2.imread(file_path), cv2.COLOR_BGR2RGB)
141 A = cv2.imread(file_path)
142 size = A.shape
143
144 # Generate the features of A.
145 print('Generating the image features...')
146 features = createDataset(np.int64(A))
147 print('Done!')
148
149 # Generate a random set of k-mean centers.
150 k = 7
151 print('\nGenerating Random k-means...')
152 k_means = generateMeans(k)
153 print(k_means)
154 print('Done!')
155
156 # Compute the k-means cluster.
157 print('\nCalculating k-mean clusters...')
158 [clusters, centers, idx] = kMeansCluster(features, k_means)
159 print('Done!')
160
161 # Display the k-mean centers.
162 print('K-mean Centers Produced After Clustering')
163 print(centers)
164
165 # Map the old image values to the new k-mean centers.
166 print('\nGenerating the new image...')
167 clustered_image = mapValues(A, idx, centers)
168 print('Done!')
169 print('')
170
171 # Show the two images.
172 plt.figure(1)
173 plt.imshow(cv2.cvtColor(np.uint8(A), cv2.COLOR_BGR2RGB))
174 plt.title('Original Image')
175 plt.xlabel('Columns')
176 plt.ylabel('Rows')
177 plt.colorbar()
178
179 plt.figure(2)
180 plt.imshow(np.uint8(clustered_image))
181 plt.title('Clustered Image')
182 plt.xlabel('Columns')
183 plt.ylabel('Rows')
184 plt.colorbar()
185 plt.show()
186

```

Output:
Input Image:



Segmented Image:



K-mean Centers:

```
K-mean Centers Produced After Clustering  
[[153 125 100]  
 [100 122 145]  
 [165 152 137]  
 [109  98  72]  
 [138 167 198]  
 [ 20  27  30]  
 [ 54  74  82]]
```

Problem 4:

Output and Questions:

1:

```
42 #-----
43 # IMPLEMENTED CONV BLOCK 2
44 self.conv_block2 = nn.Sequential(
45     nn.Conv2d(c1, c2, 3, padding=1),
46     nn.ReLU(inplace=True),
47     nn.Conv2d(c2, c2, 3, padding=1),
48     nn.ReLU(inplace=True),
49     nn.MaxPool2d(2, stride=2, ceil_mode=True),
50 )
51 #-----
52
53
54 self.conv_block3 = nn.Sequential(
55     nn.Conv2d(c2, c3, 3, padding=1),
56     nn.ReLU(inplace=True),
57     nn.Conv2d(c3, c3, 3, padding=1),
58     nn.ReLU(inplace=True),
59     nn.Conv2d(c3, c3, 3, padding=1),
60     nn.ReLU(inplace=True),
61     nn.MaxPool2d(2, stride=2, ceil_mode=True),
62 )
63
64 #-----
65 # IMPLEMENTED CONV BLOCK 4
66 self.conv_block4 = nn.Sequential(
67     nn.Conv2d(c3, c4, 3, padding=1),
68     nn.ReLU(inplace=True),
69     nn.Conv2d(c4, c4, 3, padding=1),
70     nn.ReLU(inplace=True),
71     nn.Conv2d(c4, c4, 3, padding=1),
72     nn.ReLU(inplace=True),
73     nn.MaxPool2d(2, stride=2, ceil_mode=True),
74 )
75 #-----
```


Above is the code that was written in order to complete the network model. Here, the structure of the network is a series of convolutional blocks defined as a series of convolutions followed by activations. At the end of that sequence of convolutions and activations, the activation layers are pooled.

2:

Here, we are **not** using predefined weights for our model, and we are training the model ourselves. This appears to be the case due to the structure of the *utils.py* script, in the “get_upsampling_weight” method. In this method, we can see that new weights are being calculated from incoming data.

3:

Training Loss:



Validation Loss:



4:

The metrics used by the original paper are

1. *Pixel accuracy*
2. *Mean pixel accuracy*
3. *Mean IU*
4. *Frequency-weighted IU*

Below are the results of validating the model after training:

```
Overall Acc: 0.9127122486125301
Mean Acc : 0.6298571089239537
FreqW Acc : 0.847739033039004
Mean IoU : 0.5393143651087131
0 0.9548651652445043
1 0.677273897919775
2 0.845158462670853
3 0.414655117508439
4 0.3486531279763397
5 0.24138339127602687
6 0.3013780707010186
7 0.4027415409241566
8 0.8585626209329092
9 0.5291734553865135
10 0.8817832670736081
11 0.5433616143546321
12 0.21213563688292777
13 0.8510924409864606
14 0.500840784996439
15 0.48907569632830455
16 0.4194971818451498
17 0.22319979157063552
18 0.5521416724868552
```

Here, we are measuring the

1. *Overall accuracy*
2. *Mean accuracy*
3. *Frequency-weighted accuracy*
4. *Mean IoU*

Stronger vs Weaker Classes	
<i>IoU Above 0.5</i>	<i>IoU Below 0.5</i>
Road	Wall
Sidewalk	Fence
Building	Pole
Vegetation	Traffic Light
Terrain	Traffic Sign
Sky	Rider
Person	Bus
Car	Train
Truck	Motorcycle
Bicycle	

6:

Code:

```
#-----
# Sean Carda
# ECE 253 - Image Processing
# Homework 4 - Problem 4.6
# 12/8/21
#-----

import cv2
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.colors as clr
import matplotlib.patches as ptch

# Function to plot the segmented image.
def plot(im, labels, item_list, legend_on, fignum, title):
    plt.figure(fignum, figsize=(20, 50))
    plt.imshow(cv2.cvtColor(im, cv2.COLOR_BGR2RGB))
    plt.title(title)
    plt.xlabel('Columns')
    plt.ylabel('Rows')
    if legend_on:
        plt.legend(handles=item_list)

# Read in the original image.
A_original = cv2.imread('test_image.jpg')
A_segment = cv2.imread('test_segment.jpg')


# List the colors used for segmentation.
colors = [ # [ 0, 0, 0],
    [128, 64, 128],
    [244, 35, 232],
    [70, 70, 70],
    [102, 102, 156],
    [190, 153, 153],
    [153, 153, 153],
    [250, 170, 30],
    [220, 220, 0],
    [107, 142, 35],
    [152, 251, 152],
    [0, 130, 180],
    [220, 20, 60],
    [255, 0, 0],
    [0, 0, 142],
    [0, 0, 70],
    [0, 60, 100],
    [0, 80, 100],
    [0, 0, 230],
    [119, 11, 32],
]

# List the labels.
class_names = [
    "road",
    "sidewalk",
    "building",
    "wall",
    "fence",
    "pole",
    "traffic_light",
    "traffic_sign",
    "vegetation",
    "terrain",
    "sky",
    "person",
    "rider",
    "car",
    "truck",
    "bus",
    "train",
    "motorcycle",
    "bicycle",
]

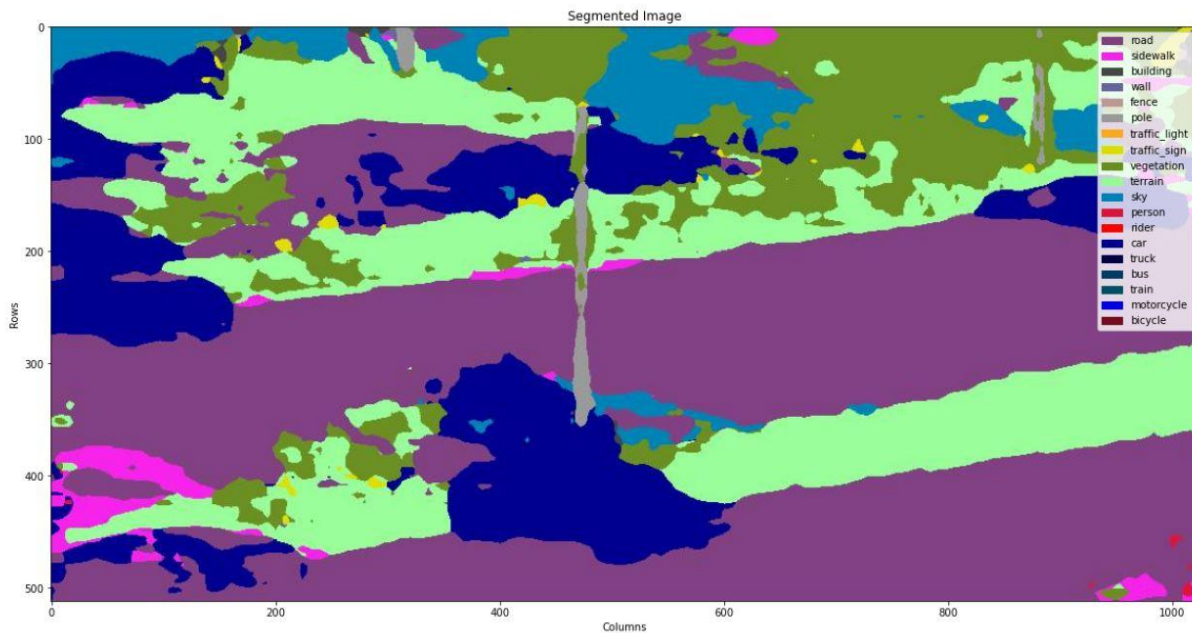
item_list = []
for i in range(len(class_names)):
    current_color = colors[i]
    current_color_tuple = (current_color[0]/255, current_color[1]/255, current_color[2]/255)
    item_list.append(ptch.Patch(color=current_color_tuple, label=class_names[i]))

plot(A_original, None, None, False, 1, 'Original Image')
plot(A_segment, class_names, item_list, True, 2, 'Segmented Image')
plt.show()
```

Output:
Original Test Image:



Segmented Classified Image:



It is difficult to determine whether or not this output looks acceptable. If we were only interested in obtaining the identity of the regions in the original image or count how many of any one object appears in the original image, then this could be sufficient. However, the segmentation seems to have a difficult time connecting the same object for a test image. For example, some of the cars are not complete, and some have “road” identified inside of them, which is clearly not the case in the original. The classifier also seems to have some difficulty in correctly identifying certain objects, such as sky (which is not contained in the original, but the

model has identified it anyway). This is interesting as “sky” was one of the labels that performed very well.

7:

Something that could have an effect, that is slightly unrelated to the actual convolutional network is perhaps segmenting the images beforehand with something like k-means clustering. With all the dynamic objects that occur in an image, having to learn so many different variations on a local level could be complicated. We saw this during training – that even after 85k iterations, the losses of the images were still immense. Narrowing the color levels to something more recognizable could lessen the level of variety required. An adverse effect of this is that the number of objects that can be classified might be less depending on the value of k used (if using k-means clustering).

A simpler option could be to utilize fewer class labels. Using fewer labels might reduce the training error by giving the model less to identify. As we could see from the validation, there were several labels that were weakly identified by the model. If we threw those out, we might have a better time predicting the labels that were often identified by the model.

Problem 5:

Code:

```
1 #
2 # Sean Carda
3 # ECE 253 - Image Processing
4 # Dr. Mohan Trivedi
5 # Homework 3
6 #-----
7
8 #-----
9 # Global Imports.
10 import cv2
11 import numpy as np
12 import scipy.ndimage as ndi
13 import matplotlib.pyplot as plt
14 #-----
15 # Suppress scientific notation representation for numpy arrays.
16 np.set_printoptions(suppress=True)
17
18 # -----
19 # This filter will load an image, perform canny edge operation on it, and replace the white
20 # pixels with colored pixels.
21 #-----
22 file_path = 'lane.png'
23 A = cv2.imread(file_path)
24
25 # Compute the canny edge of the image.
26 A_edge = 255 * cv2.Canny(A, 100, 200, i2gradient=True)
27 edge_size = A_edge.shape
28
29 # Layer the edged image.
30 output = np.zeros([edge_size[0], edge_size[1], 3])
31 output[:, :, 0] = A_edge
32 output[:, :, 1] = A_edge
33 output[:, :, 2] = A_edge
34
35 # For every columns, replace it with a colored pixel column.
36 for i in range(edge_size[1]):
37     new_pixel = np.random.randint(1, 255, size=(1, 3))
38     output[:, i, :] = np.where(output[:, i, :] > 0, new_pixel, 0)
39
40 # Show the image.
41 plt.figure(1)
42 plt.imshow(np.uint8(output))
43 plt.title('Colorful Lanes!')
44 plt.xlabel('Columns')
45 plt.ylabel('Rows')
46 plt.colorbar()
47 plt.show()
```

Output:

